

Verteilte Systeme/ Distributed Systems

Artur Andrzejak

1


Motivation and Introduction

Distributed System (DS) - Definitions

- ▶ A distributed system is one in which the failure of a computer you didn't even know existed can render your own computer unusable.”

Leslie Lamport

- ▶ [Col] A distributed system is one in which components..
 - ▶ are located at networked computers
 - ▶ communicate and coordinate only by passing messages



More Definitions

- ▶ A distributed computing system
 - ▶ consists of multiple autonomous processors
 - ▶ that do not share primary memory
 - ▶ but cooperate by sending messages over a communication network
- ▶ Consequences
 - ▶ concurrency of actions
 - ▶ lack of a global clock
 - ▶ independent failures
- ▶ A lot of material in this lecture is devoted to dealing with these issues!

Henri E. Bal

Common Challenges in Distributed Systems

▶ **Higher complexity of software**

- ▶ Algorithms, implementation and debugging much more involved than in „serialized“ software
- ▶ A **BIG** problem now, especially with multi-core CPUs

▶ **Dependency on the underlying network**

- ▶ Differences in message transmission times
- ▶ Non-deterministic phenomena, e.g. order of message arrival

▶ **Higher failure rate, lower reliability**

- ▶ My internet connection at home (cable) malfunctions at least once a month
 - ▶ A generic laptop breaks down every 2-4 years

Other Challenges in Distributed Systems

1. Security

- ▶ Privacy
- ▶ Authentication
- ▶ Availability

2. Scalability

- ▶ e.g. IP addresses become scarce: in IPv4, we have only 2^{32} ~ $4 \cdot 10^9$ IP addresses
- ▶ Ongoing effort: from 32 to 128 bits in IPv6

3. Heterogeneity of ...

- ▶ network infrastructure
- ▶ computer hard- and software (e.g., operating systems)
- ▶ data representation in protocols

4. Concurrency

- ▶ Avoidance of deadlocks and race conditions


Why are Distributed Systems Needed at all?

- ▶ [Col] “Sharing of resources is a main motivation for constructing distributed systems”
- ▶ „Sharing of resources“ has two reasons/motivations:
 - ▶ R1. Distributing computations over several cores/CPU's
 - ▶ R2. Information/data sharing and communication
- ▶ Reason 1 leads to distributed / parallel computing
 - ▶ Related to scientific computing / number crunching and non-consumer applications (i.e. science, simulations)
- ▶ Reason 2 is related to end-users/“business”
 - ▶ Most consumers buy computers for data/information exchange (email, WWW), not for “compute jobs”


Reason 1: Why distributed/parallel computing?

- ▶ **Cumulative power**: distributed systems consisting of collections of microcomputers may have processing powers that no single computer will ever achieve
- ▶ **Economics**: collections of microprocessors offer a better price/performance ratio than mainframes
 - ▶ Why is Google using (hundred of thousands) of commodity-like PCs?
- ▶ **CPU architectures**: we experience a shift in CPU-architectures from “one core, higher frequency”, to “many cores, same frequency”

Reason 1: Moore's Law and CPU architectures



- ▶ Moore's Law: „*The density of transistors on a chip doubles every 18 months, for the same cost (1965)*“



- ▶ Since about 2006, performance of a single core does not increase at this rate
- ▶ But Moore's Law still valid => more cores per chip
- ▶ Parallelization is needed to increase computing power

Reason 1: Supercomputers Become Massively Parallel

- ▶ Parallelization is the primary way to speed up computations
- ▶ Processor count in supercomputers increases each year
 - ▶ Speed of a rank X supercomputer doubles every 13 months
 - ▶ But # transistors ~ # cores per chip doubles only every 24 months
- ▶ This implies higher costs - mainly due to el. energy

	Cray X-MP/24	Cray Y-MP4/264	Cray T3D SC 192	Cray T3E	IBM 690 pSeries	SGI ICE
Year	1986	1992	1994	1997	2002	2008
GFlops/s	0.47	1.33	38	363	2662	125,000
# CPUs	2	4	192	512	512	~2500
Power (kW)				90	160	600




Reason 2 for Distributed Systems

- ▶ Information/data sharing, and communication
- ▶ WWW as a main distributed application

Date	Computers	Web servers	Percentage
1993, July	1,776,000	130	0.008
1995, July	6,642,000	23,500	0.4
1997, July	19,540,000	1,203,096	6
1999, July	56,218,000	6,598,697	12
2001, July	125,888,197	31,299,592	25
2003, January	171,638,297	35,424,956	20

Illustration of Reason 2: Communication and Sharing of Resources in Intranet



System Architectures

Taxonomy of Architectures

- ▶ Introduced 1972 by Michael Flynn

		<i>number of instruction streams</i>	
		single	multiple
<i>number of data streams</i>	single	SISD – single-threaded process	MISD – pipeline architecture (rare)
	multiple	SIMD – vector processing	MIMD – multi-threaded programming

Michael Flynn's Taxonomy of Architectures

*number of **instruction** streams*

single

multiple

*number of **data** streams*

single

SISD

- Traditional one-core system

MISD

- Generally not used and doesn't make sense, except to classifying redundant systems (tandem-computers)

SIMD

- Array (vector) processor, e.g.
 - GPUs
 - SSE3: Intel's Streaming SIMD Extensions


MIMD

- Multiple computers / cores, each with:
 - program counter, program, data
- Both parallel and distributed systems

multiple

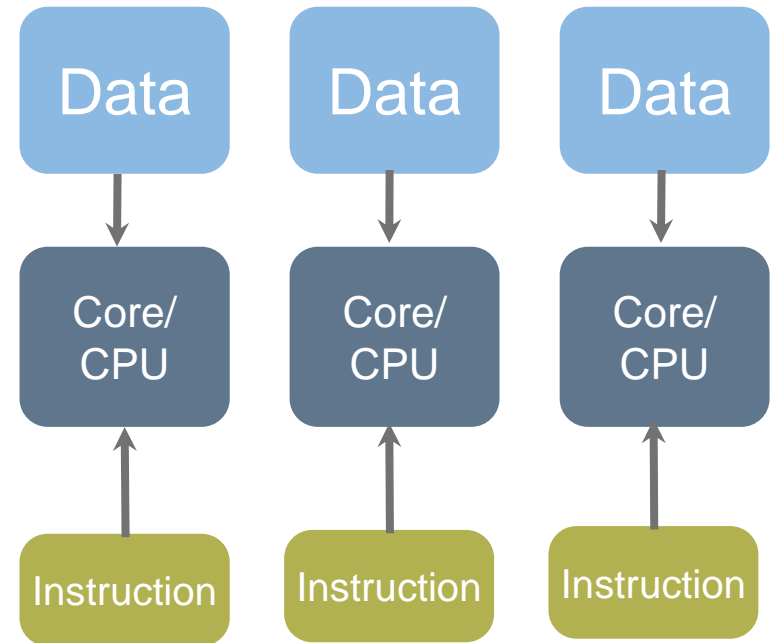
Flynn's Taxonomy: SIMD vs. MIMD

SIMD



E.g. GPUs, MMX, SSE2, vector computers

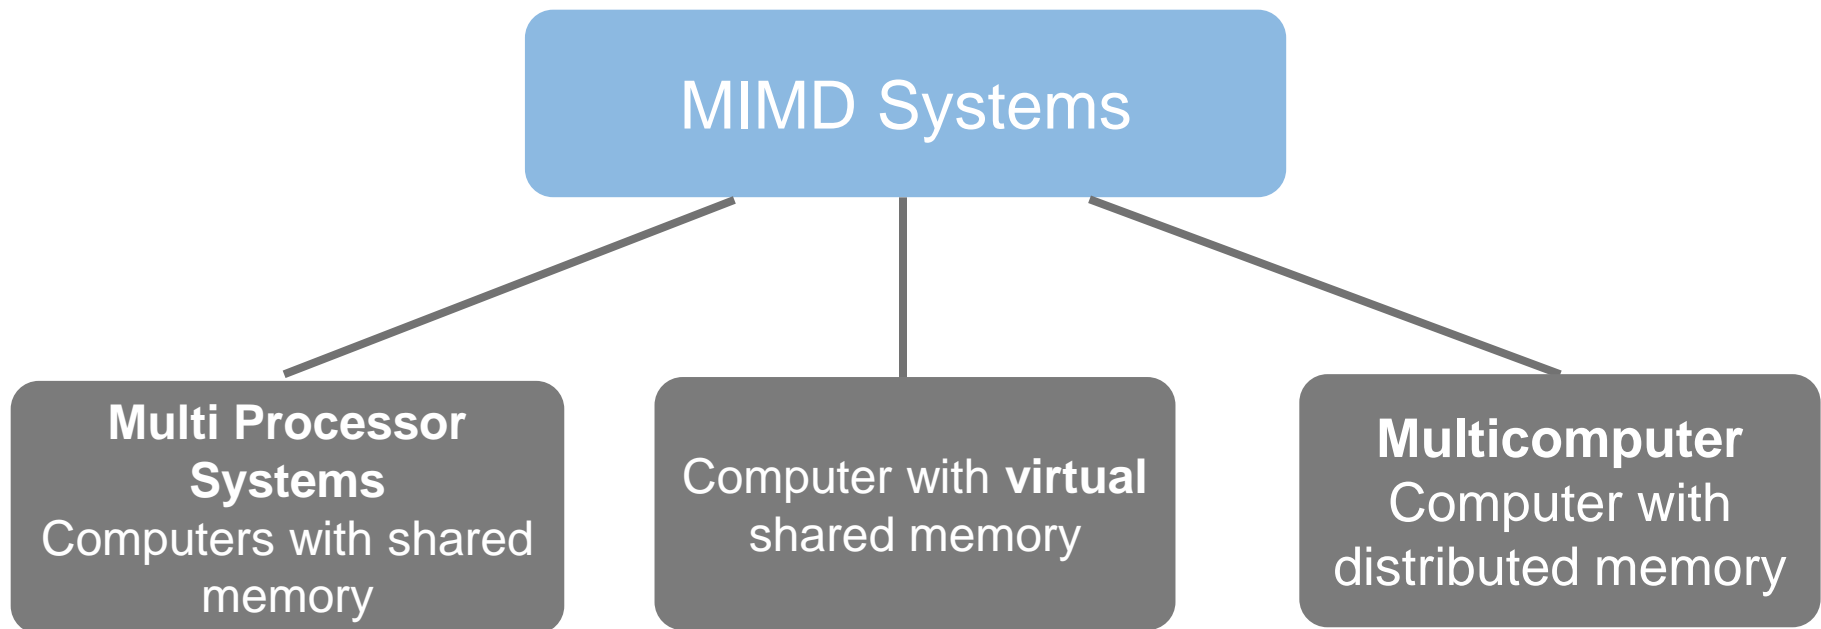
MIMD



E.g. power workstation, cluster, supercomputer

Memory Organization

- ▶ Most parallel computers have MIMD-architecture
- ▶ However, they can have various memory organization / architecture:



Parallel and Distributed Computing Systems

▶ **Parallel computing** systems

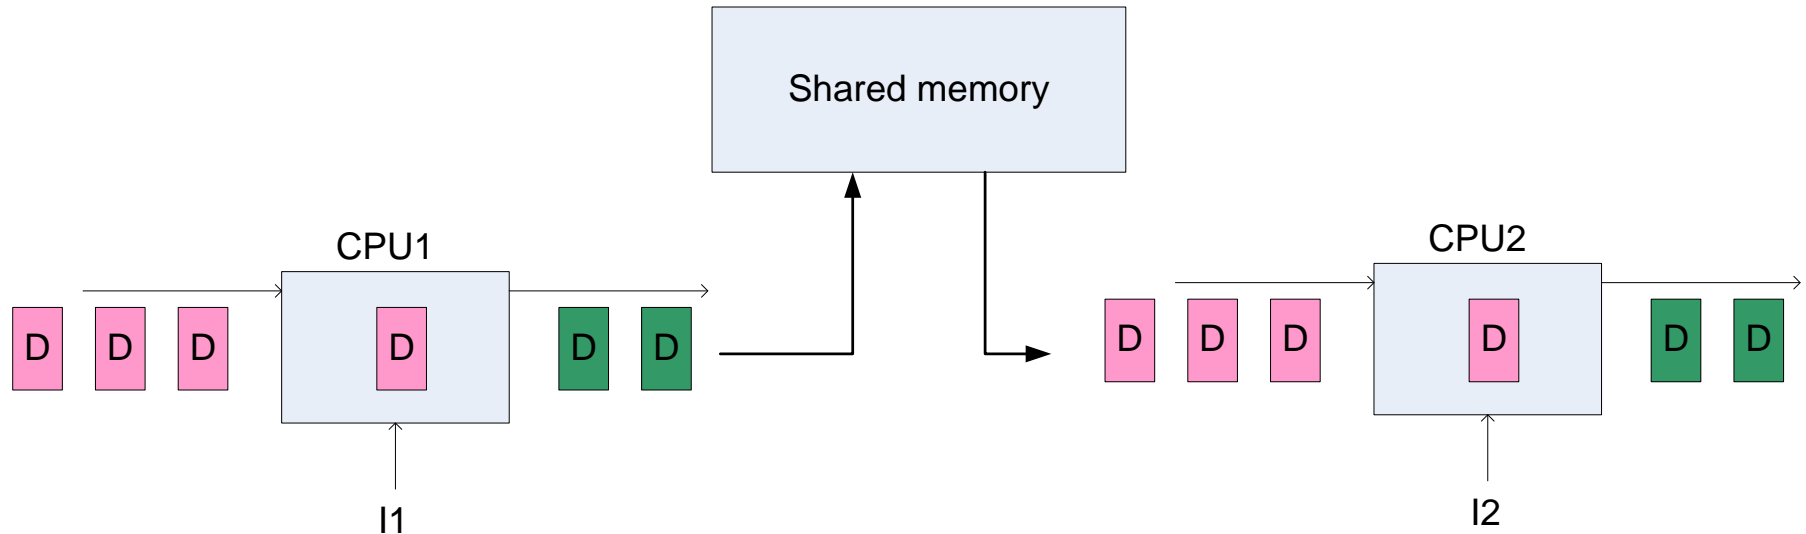
- ▶ Multi- and many-core processors
- ▶ **Multiprocessor** systems
 - ▶ Multiple CPUs in the same „box“
- ▶ Vector processing of data (e.g. Cray-Systems)

▶ **Distributed computing** systems

- ▶ **Multicomputers**: multiple CPUs across many computers (independent nodes)
- ▶ Communication via (high-speed) network

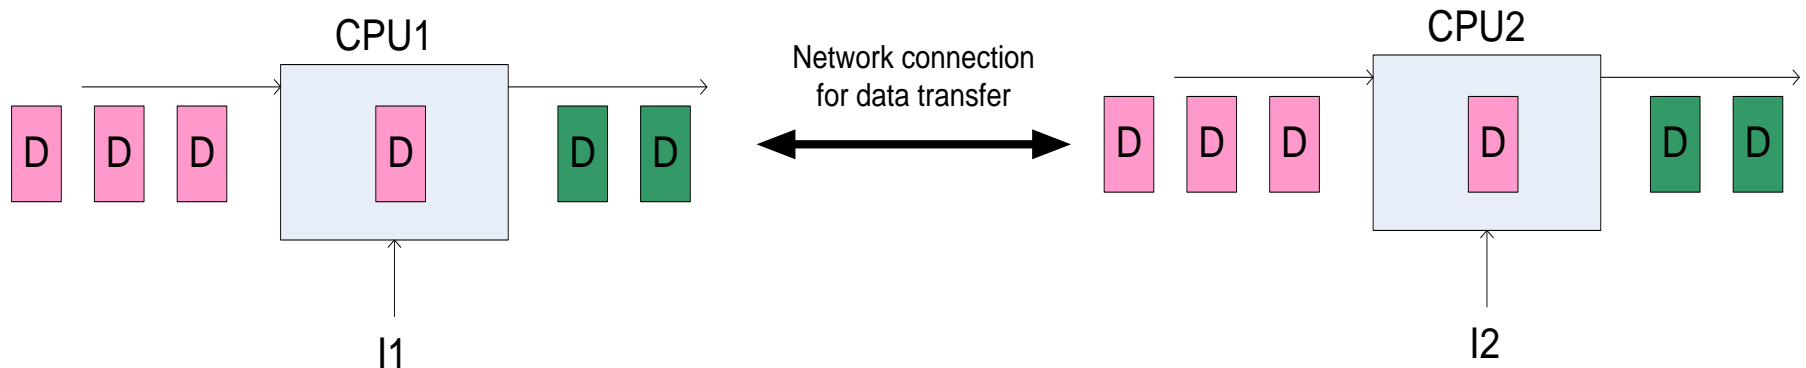
Parallel, but not Distributed Hardware

- ▶ **Multi processor machines** perform parallel processing via multiple CPUs communicating over the same **shared-memory** hardware/bus



Multicomputers: „Truly Distributed“ Hardware

- ▶ A **multicomputer** has following architecture:
 - ▶ It has multiple independent computers / nodes
 - ▶ Each node has its own memory
 - ▶ They communicate over a (private or public) network
 - ▶ Usually via message passing, e.g. MPI



Communication Styles 1: Hardware vs. Communication

- ▶ DS definitions say that it does not share primary memory and nodes communicate only by message passing
- ▶ In fact, these are two independent aspects
 - ▶ A. Memory organization: shared vs. not shared memory
 - ▶ B. Communication type: “shared mem” vs. “message passing”

	A1. “True” single bus/memory hardware (PCs, workstations etc.)	A2. Physically distributed nodes (clusters, clouds, HPC systems, ...)
B1. Communication via shared-memory	“Natural”, but coding is error-prone (race conditions, deadlocks..)	Possible as “Distributed Shared Memory”; lot of research, controversial
B2. Communication via message passing	Used in “safe” OSs and modern prg. languages: Actors in Erlang, Scala, ..	“Natural” and widely used, e.g. MPI; can be more abstract, e.g. Java RMI

Communication Styles 2: Shared-Memory vs. Message Passing

- ▶ Both form of communication are used in distributed and non-distributed systems!
- ▶ **Distributed Shared-Memory**
 - ▶ A communication approach, in which each node of a distributed system has access to a large shared memory (in addition to own non-shared private memory)
- ▶ **Actor model** of concurrent programming
 - ▶ A model of concurrent processing where threads communicate via message passing and do not use locks / monitors; i.e. processors with physically shared memory still use message passing
 - ▶ Available in Erlang, Scala, Java libraries, some OSs..

Thank you.