

Verteilte Systeme/ Distributed Systems

Artur Andrzejak

2

Overview

Interprocess Communication (IPC)

Overview IPC

- ▶ Within an Operating System (**OS**) threads and processes need to communicate
 - ▶ Threads can use shared variables to communicate
- ▶ For processes it is harder: they do not share memory
 - ▶ Local processes: different address space / protection
 - ▶ Remote processes: physical barrier
- ▶ Consequently, a set of techniques for the exchange of data among processes exist – they are called **Inter-process communication (IPC)** methods
- ▶ We discuss major ones, more exotic in appendix



IPC Types [local]: Files, Semaphores, Signals

- ▶ **Files:** all operating systems
- ▶ **Semaphores:** all POSIX systems; Windows
 - ▶ Synchronizes access to shared data and resources
 - ▶ A protected variable or abstract data type
 - ▶ Only one thread / process can change it at a time
- ▶ **Signals:** most OSs; Windows: only C run-time lib
 - ▶ Essentially it is an asynchronous notification sent to a process in order to notify it of an event that occurred
 - ▶ When a signal is sent to a process, the operating system interrupts the process's normal flow of execution



IPC Types [local]: Shared memory, Memory-mapped Files

- ▶ **Shared Memory:** all POSIX systems
 - ▶ Memory that may be simultaneously accessed by multiple programs with an intent to provide communication among them or avoid redundant copies
 - ▶ Used e.g. to transfer images between the application and the X server on Unix systems
- ▶ **Memory-mapped files:** POSIX; Windows ([Link](#))
 - ▶ **mmap** is a POSIX-compliant Unix system call that maps files or devices into memory
 - ▶ Initially file contents are not entirely read from disk and don't use physical RAM at all
 - ▶ The actual reads from disk are performed in "lazy" manner, after a specific location is accessed ([demand paging](#))



IPC Types [local/remote]: Pipes

- ▶ **Pipes** (anonymous/named): all POSIX systems
 - ▶ A pipe is a buffered data stream between two processes
 - ▶ FIFO; treated as files on API-level
 - ▶ a pipe(line): a set of processes chained by their standard streams, so that the output of each process (stdout) feeds directly as input (stdin) of the next one
- ▶ Example in a POSIX-shell
 - ▶ `ps | grep java`
 - ▶ Output of “ps” is piped to a line filter, only lines containing “java” are printed on stdout
- ▶ Programmatically: **pipe** (int fd[])
 - ▶ fd[0] is a reference (handle) to “read-from” end (pipe output)
 - ▶ fd[1] is a reference (handle) to “write-to” end (pipe input)



IPC Types [local/remote]: Sockets, MP

- ▶ **Sockets:** most OSs

- ▶ An API which allows communications between hosts or between processes on one computer
- ▶ Using the concept of an Internet socket
- ▶ More on this today

- ▶ **Message passing:**

- ▶ Via Middleware: mainly **MPI**, Message Passing Interface
- ▶ Communication via sending of messages to recipients
- ▶ (Much) more on this later



IPC Types [local/remote]: RPC

▶ **RPC – Remote Procedure Calls**

- ▶ It provides an API for calling a “normal” function
 - ▶ However, this function can be in the same process, or another process or even host (= computer)
 - ▶ Call parameters and return values are the actual data to be exchanged
 - ▶ Greatly simplifies programming of distributed applications
-
- ▶ Much more on this will follow:
 - ▶ General RPC concepts
 - ▶ Java RMI
 - ▶ Google RPC (gRPC) and protocol buffers
 - ▶ Web services, esp. REST

Sockets: Introduction

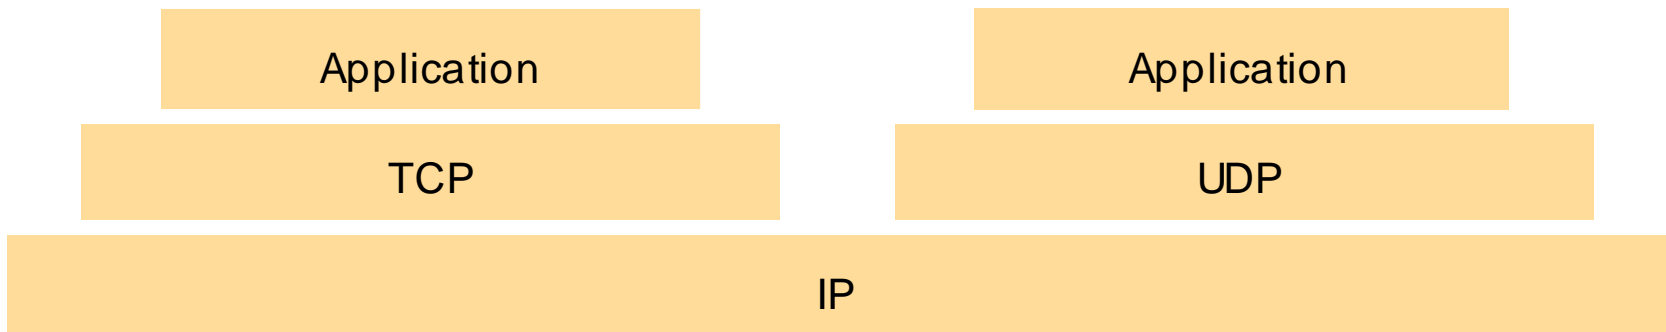
IP – Internet protocol

- ▶ Born 1969 as a research network of 4 machines
- ▶ Funded by Department of Defense's ARPA
 - ▶ Advanced Research Projects Agency, renamed to Defense Advanced Research Projects Agency (DARPA) in 1972
- ▶ In early 1970s used in ARPANET, first large-scale computer network
- ▶ Together with the TCP layer builds **TCP/IP protocol stack**
- ▶ Goal
 - ▶ *Build an efficient fault-tolerant network that could connect heterogeneous machines and link separately connected networks*

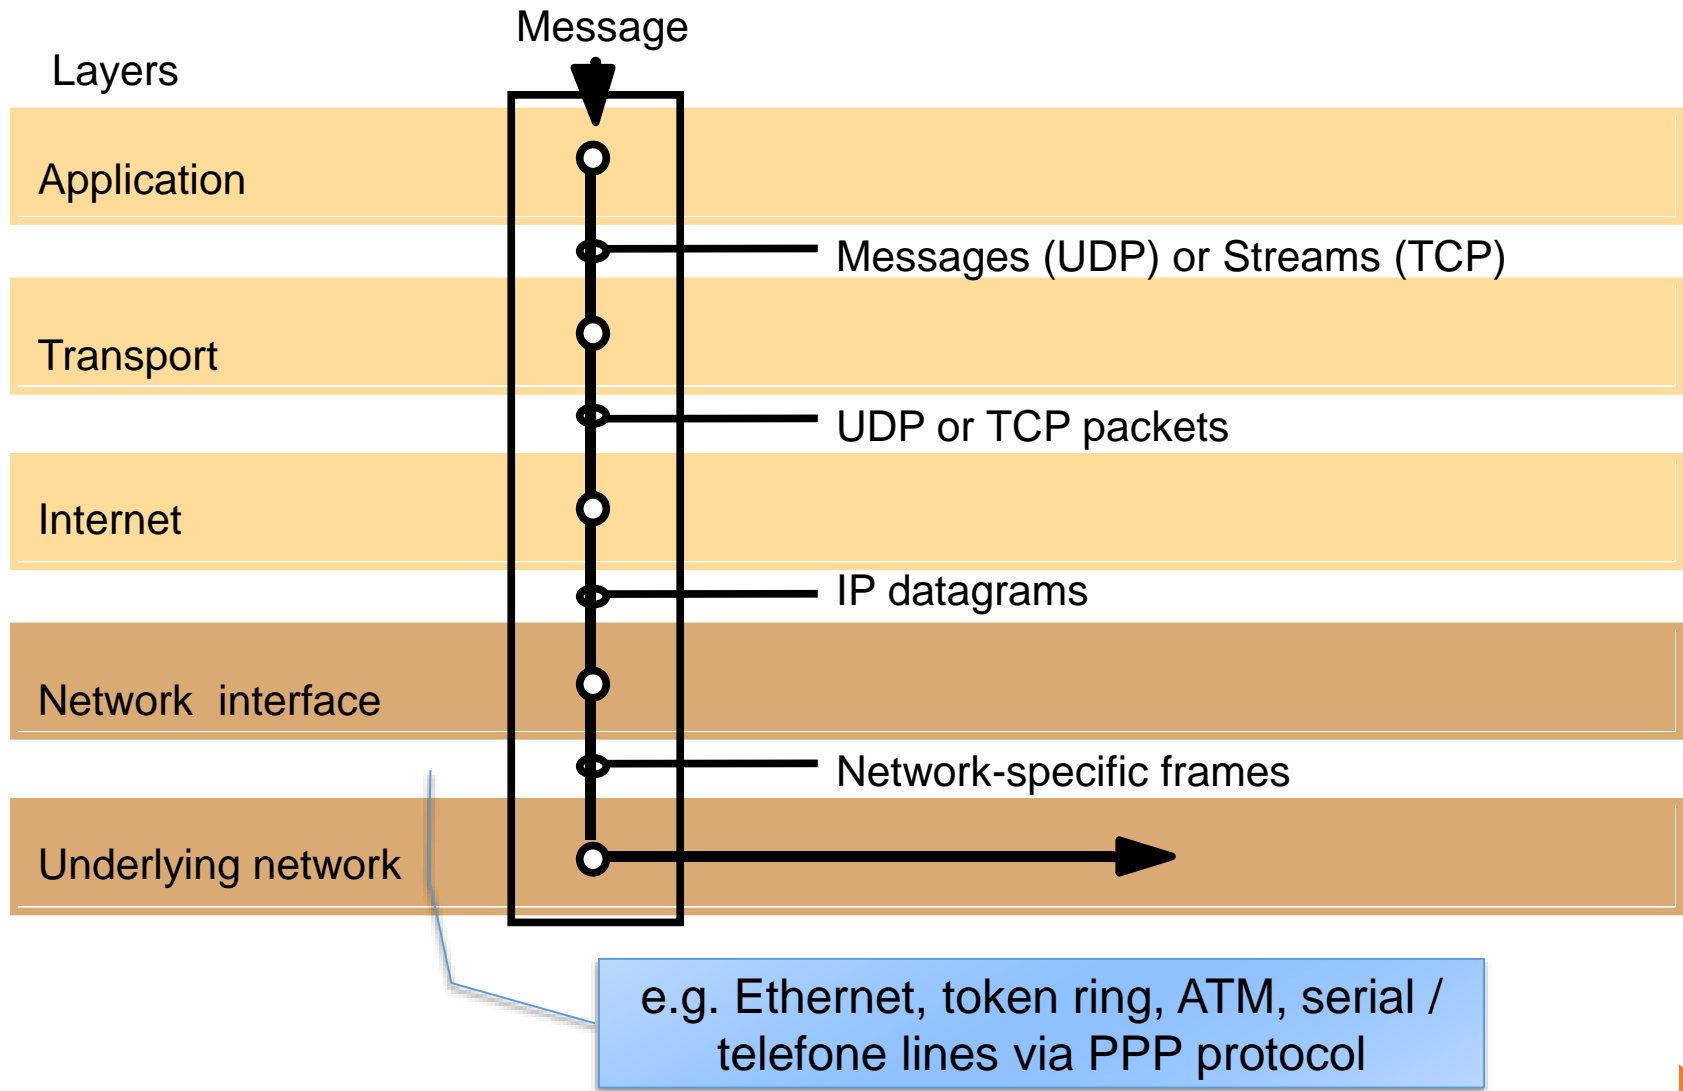


Layers of the Internet protocol stack

- ▶ **IP (Internet Protocol)** is the underlying network protocol of the Internet virtual network
 - ▶ Logical network on top of multiple physical networks
 - ▶ Provides basic transmission mechanisms
- ▶ Based on it are two transport protocols
 - ▶ **TCP: Transport Control Protocol**
 - ▶ Reliable connection-oriented protocol
 - ▶ **UDP: User Datagram Protocol**
 - ▶ Datagram protocol - does not guarantee reliable transmission



TCP/IP layers



Socket

- ▶ **Socket**: an application-created, OS-controlled interface/API (a “door”) into which application process can both send and receive messages to/from another application process
- ▶ Sockets are easy to use: they mimic file operations
 - ▶ connect
 - ▶ read/write (stream) or send/receive (datagram)
 - ▶ disconnect
- ▶ **Socket API**
 - ▶ introduced in BSD4.1 UNIX, 1981
 - ▶ explicitly created, used, released by apps
 - ▶ client/server paradigm



Sockets and Ports

▶ Port

- ▶ a **message destination** specified by a small integer (16 bits)
- ▶ Any process can send a message to it
- ▶ Internet protocols use the combination (IP address, local port)
- ▶ IANA (Internet Assigned Numbers Authority) ports:
 - ▶ well-known ports: 1 – 1023
 - ▶ registered ports: 1024 – 49151
 - ▶ dynamic or private ports: 49152 - 65535

▶ Ports and sockets

- ▶ A socket must be bound to a local port
- ▶ A socket pair uniquely identifies a communication session
 - ▶ Pair: [local IP address, local port] + [remote IP address, remote port]



Types of Communication

- ▶ Sockets support both types of communication
 - ▶ **connectionless**
 - ▶ Via **datagrams** - implemented on network layer via **UDP**
 - ▶ **connection-oriented**
 - ▶ **Stream** communication - implemented on network layer via **TCP**
- ▶ Type must be specified when creating socket
 - ▶ Uses different calls for data exchange
- ▶ Berkeley Sockets:
 - ▶ Connectionless: **sendto**, **sendmsg**; **recvfrom**, **recvmsg**
 - ▶ Connection-oriented: **read** / **write**; **recv** / **send** (extra flags)



Sockets Programming (Mainly with Python)

Socket programming basics

- ▶ Server must be running before client can send anything to it
- ▶ Server must have a socket through which it receives and sends segments
- ▶ Similarly client needs a socket
- ▶ Socket is locally identified with a port number
- ▶ Client needs to know server IP address and socket port number

Socket Programming with *UDP*

- ▶ **UDP: no “connection” between client and server**
- ▶ no handshaking
- ▶ sender explicitly attaches IP address and port of destination to each segment
- ▶ OS attaches IP address and port of sending socket to each segment
- ▶ Server can extract IP address, port of sender from received segment
- ▶ UDP provides unreliable transfer of groups of bytes (“datagrams”) between client and server
- ▶ Programming with UDP is “straightforward”
 - ▶ See example of Java client/server in additional slides

Datagram (UDP) Communication in C

Sending a message (client)

```
s = socket(AF_INET, SOCK_DGRAM, 0)
•
•
bind(s, ClientAddress)
•
•
sendto(s, "message", ServerAddress)
```

Receiving a message (server)

```
s = socket(AF_INET, SOCK_DGRAM, 0)
•
•
bind(s, ServerAddress)
•
•
amount = recvfrom(s, buffer, from)
```

socket()

AF_INET – communication domain is Internet domain

SOCK_DGRAM – specifies datagram communication

0 – lets the system select the protocol (here UDP)

bind()

ServerAddress and *ClientAddress* are socket addresses

sender: uses any available local port

receiver: uses a “well-known” port of itself

Supplies sender's address with each message it delivers



Socket Programming with *TCP*

Client must contact server

- ▶ server process must first be running
- ▶ server must have created socket that welcomes client's contact

Client contacts server by:

- ▶ creating client-local TCP socket
- ▶ specifying IP address, port number of server process
- ▶ When *client creates socket*: client TCP establishes connection to server TCP

- ▶ When contacted by client, *server TCP creates new socket* for server process to communicate with client
 - ▶ allows server to talk with multiple clients
 - ▶ source port numbers used to distinguish clients

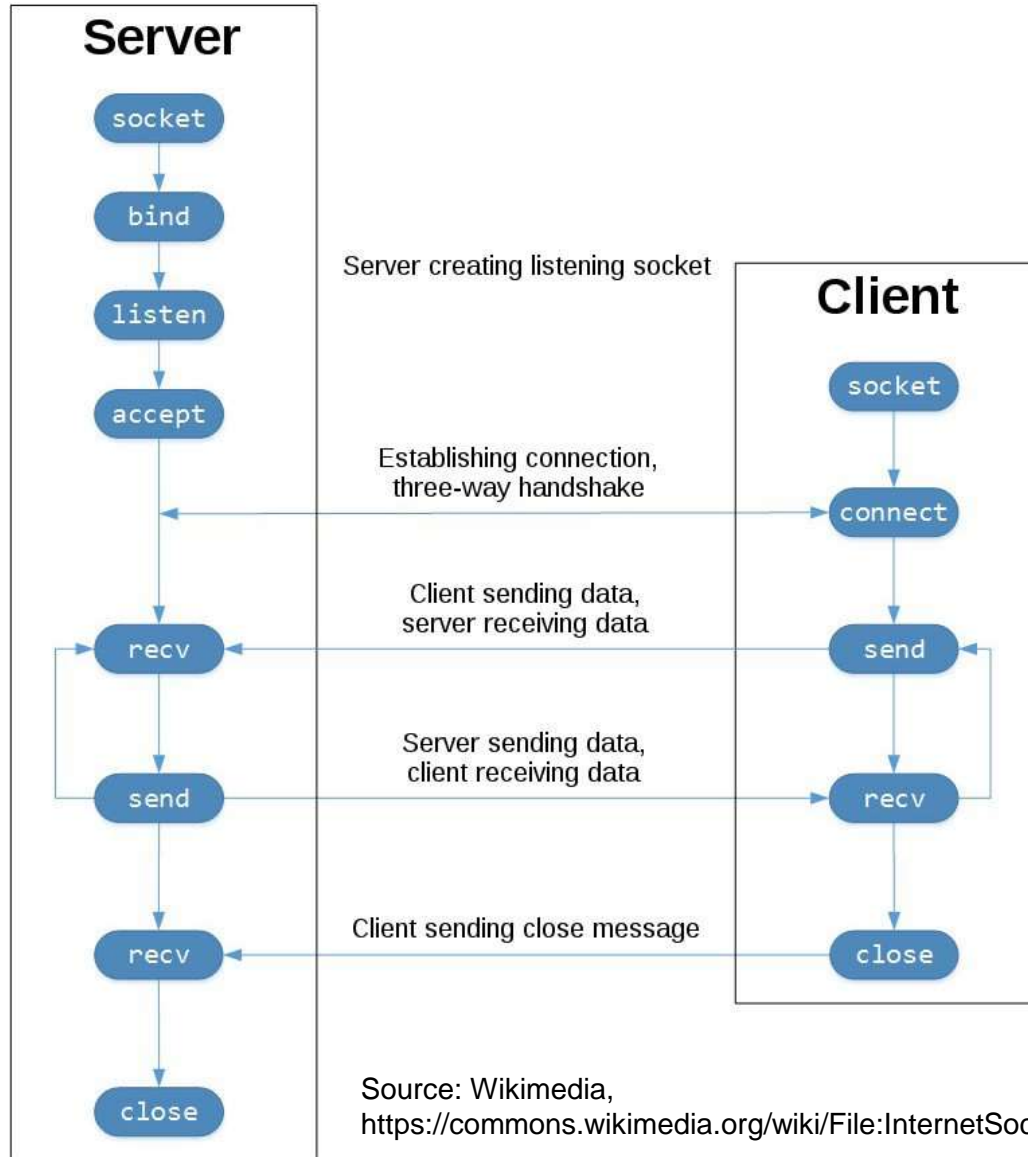
application viewpoint

TCP provides reliable, in-order transfer of bytes ("pipe") between client and server

TCP-Sockets in Python

- ▶ Based on the tutorial **Socket Programming in Python** by Nathan Jennings (<https://realpython.com/python-sockets/>)
- ▶ We use Python's [socket module](#) as an interface to the Berkeley sockets API
 - ▶ Remote hosts are accessible
 - ▶ Compare to: [Unix domain sockets](#) - allow only same-host communication
- ▶ Higher-level APIs available in Python:
 - ▶ [socketserver](#) module, a framework for network servers
 - ▶ Support for various protocols ([link](#)):
 - ▶ HTTP, FTP, POP, IMAP, SMTP, telnet, XMLRPC, ...

Overview of the Interaction



TCP Server in Python

```
def prepareReply(inData):
```

```
    #...
```

```
import socket
```

```
HOST = "127.0.0.1"
```

```
PORT = 65432
```

```
with socket.socket (socket.AF_INET, socket.SOCK_STREAM) as s:
```

```
    s.bind((HOST, PORT))
```

```
    s.listen()
```

```
    conn, addr = s.accept()
```

```
    with conn:
```

```
        print("Connected by", addr)
```

```
        while True:
```

```
            data = conn.recv(1024)
```

```
            if not data:
```

```
                break
```

```
            response = prepareReply(data)
```

```
            conn.sendall(response)
```

TCP Server in Python

My “secret” function

```
def prepareReply(inData):
```

```
    #...
```

```
import socket
```

```
HOST = "127.0.0.1"
```

```
PORT = 65432
```

Creates a socket object that supports the **context manager type**, so there is no need to call `s.close()`

```
with socket.socket (socket.AF_INET, socket.SOCK_STREAM) as s:
```

```
    s.bind((HOST, PORT))
```

```
    s.listen()
```

```
    conn, addr = s.accept()
```

```
    with conn:
```

```
        print("Connected by", addr)
```

```
        while True:
```

```
            data = conn.recv(1024)
```

```
            if not data:
```

```
                break
```

```
            response = prepareReply(data)
```

```
            conn.sendall(response)
```


TCP Server in Python

```
def prepareReply(inData):
```

```
    #...
```

```
import socket
```

```
HOST = "127.0.0.1"
```

```
PORT = 65432
```

Specify the address family and socket type:

- AF_INET is the Internet address family for IPv4
- SOCK_STREAM is the socket type for TCP

```
with socket.socket(socket.AF_INET, socket.SOCK_STREAM) as s:
```

```
    s.bind((HOST, PORT))
```

```
    s.listen()
```

```
    conn, addr = s.accept()
```

```
    with conn:
```

```
        print("Connected by", addr)
```

```
        while True:
```

```
            data = conn.recv(1024)
```

```
            if not data:
```

```
                break
```

```
            response = prepareReply(data)
```

```
            conn.sendall(response)
```

TCP Server in Python

```
def prepareReply(inData):
```

```
    #...
```

```
import socket
```

```
HOST = "127.0.0.1"
```

```
PORT = 65432
```

```
with socket.socket(socket.AF_INET, socket.SOCK_STREAM) as s:
```

```
    s.bind((HOST, PORT))
```

```
    s.listen()
```

```
    conn, addr = s.accept()
```

```
    with conn:
```

```
        print("Connected by", addr)
```

```
        while True:
```

```
            data = conn.recv(1024)
```

```
            if not data:
```

```
                break
```

```
            response = prepareReply(data)
```

```
            conn.sendall(response)
```

Associate the socket with a specific network interface (here HOST = 127.0.0.1) and port number (here PORT = 65432)

- HOST can be a hostname, IP address, or empty string
- For empty string, the server will accept connections on all available IPv4 interfaces

TCP Server in Python

```
def prepareReply(inData):
```

```
    #...
```

```
import socket
```

```
HOST = "127.0.0.1"
```

```
PORT = 65432
```

```
with socket.socket(socket.AF_INET, socket.SOCK_STREAM) as s:
```

```
    s.bind((HOST, PORT))
```

```
    s.listen()
```

```
    conn, addr = s.accept()
```

```
    with conn:
```

```
        print("Connected by", addr)
```

```
        while True:
```

```
            data = conn.recv(1024)
```

```
            if not data:
```

```
                break
```

```
            response = prepareReply(data)
```

```
            conn.sendall(response)
```

Enables a server to accept() connections:
it makes it a “listening” socket

- **listen()** has a *backlog* parameter
- It specifies the number of unaccepted connections that the system will allow before refusing new connections

TCP Server in Python

```
def prepareReply(inData):  
    #...
```

```
import socket  
HOST = "127.0.0.1"  
PORT = 65432
```

Important: we now have a **new** socket object `conn` used to communicate with the client. It's distinct from the listening socket that the server is using to accept new connections!

```
with socket.socket (socket.AF_INET, socket.SOCK_STREAM) as s:
```

```
    s.bind((HOST, PORT))  
    s.listen()
```

```
    conn, addr = s.accept()
```

```
    with conn:
```

```
        print("Connected by", addr)
```

```
        while True:
```

```
            data = conn.recv(1024)
```

```
            if not data:
```

```
                break
```

```
            response = prepareReply(data)
```

```
            conn.sendall(response)
```

- Blocks and waits for an incoming connection
- When a client connects, it returns a **new socket object** representing the connection (`conn`) and a tuple holding the address of the client (`addr`)
- `addr` will contain (host, port) for IPv4 connections

TCP Server in Python

```
def prepareReply(inData):
```

```
    #...
```

```
import socket
```

```
HOST = "127.0.0.1"
```

```
PORT = 65432
```

```
with socket.socket(socket.AF_INET, socket.SOCK_STREAM) as s:
```

```
    s.bind((HOST, PORT))
```

```
    s.listen()
```

```
    conn, addr = s.accept()
```

```
    with conn:
```

```
        print("Connected by", addr)
```

```
        while True:
```

```
            data = conn.recv(1024)
```

```
            if not data:
```

```
                break
```

```
            response = prepareReply(data)
```

```
            conn.sendall(response)
```

- An infinite while loop is used to loop over blocking calls to `conn.recv()`.
- This reads whatever data the client sends us
- If `conn.recv()` returns an empty bytes object, then the client closed the connection and the loop is terminated

TCP Server in Python

```
def prepareReply(inData):
```

```
    #...
```

```
import socket
```

```
HOST = "127.0.0.1"
```

```
PORT = 65432
```

```
with socket.socket(socket.AF_INET
```

```
    s.bind((HOST, PORT))
```

```
    s.listen()
```

```
    conn, addr = s.accept()
```

```
    with conn:
```

```
        print("Connected by", addr)
```

```
        while True:
```

```
            data = conn.recv(1024)
```

```
            if not data:
```

```
                break
```

```
            response = prepareReply(data)
```

```
            conn.sendall(response)
```

- An infinite while loop is used to loop over blocking calls to `conn.recv()`
- This reads whatever data the client sends us
- If `conn.recv()` returns an empty bytes object, then the client closed the connection and the loop is terminated
- The **with** statement automatically closes the socket at the end of the block

TCP Server in Python

```
def prepareReply(inData):
```

```
    #...
```

```
import socket
```

```
HOST = "127.0.0.1"
```

```
PORT = 65432
```

```
with socket.socket (socket.AF_INET, socket.SOCK_STREAM) as s:
```

```
    s.bind((HOST, PORT))
```

```
    s.listen()
```

```
    conn, addr = s.accept()
```

```
    with conn:
```

```
        print("Connected by", addr)
```

```
        while True:
```

```
            data = conn.recv(1024)
```

```
            if not data:
```

```
                break
```

```
            response = prepareReply(data)
```

```
            conn.sendall(response)
```

- The received data is used to generate a reply (via my “secret” function)
- Response is sent as one “block” using **sendall()**

TCP Client in Python

```
import socket, sys  
HOST = "127.0.0.1"  
PORT = 65432
```

- Read given command line argument

```
inputStr = sys.argv[1]  
with socket.socket (socket.AF_INET, socket.SOCK_STREAM) as s:  
    s.connect((HOST, PORT))  
    s.sendall(inputStr.encode())  
    data = s.recv(1024)  
  
print("Received", repr(data.decode()))
```


TCP Client in Python

```
import socket, sys  
HOST = "127.0.0.1"  
PORT = 65432
```

- Create a TCP socket
- Same arguments as for the server

```
inputStr = sys.argv[1]  
with socket.socket (socket.AF_INET, socket.SOCK_STREAM) as s:  
    s.connect((HOST, PORT))  
    s.sendall(inputStr.encode())  
    data = s.recv(1024)  
  
print("Received", repr(data.decode()))
```

TCP Client in Python

```
import socket, sys
HOST = "127.0.0.1"
PORT = 65432
```

- Connects to the server (3-way handshake)
- Sends the cmd line string (as a byte-like object) to the server

```
inputStr = sys.argv[1]
with socket.socket (socket.AF_INET, socket.SOCK_STREAM) as s:
    s.connect((HOST, PORT))
    s.sendall(inputStr.encode())
    data = s.recv(1024)

print("Received", repr(data.decode()))
```

TCP Client in Python

```
import socket, sys
HOST = "127.0.0.1"
PORT = 65432
```

```
inputStr = sys.argv[1]
```

```
with socket.socket(socket.AF_INET, socket.SOCK_STREAM) as s:
```

```
    s.connect((HOST, PORT))
```

```
    s.sendall(inputStr.encode())
```

```
    data = s.recv(1024)
```

- Receive data from the server
- Object **socket** is closed automatically

```
print("Received", repr(data.decode()))
```

- Converts the received data to string and prints it

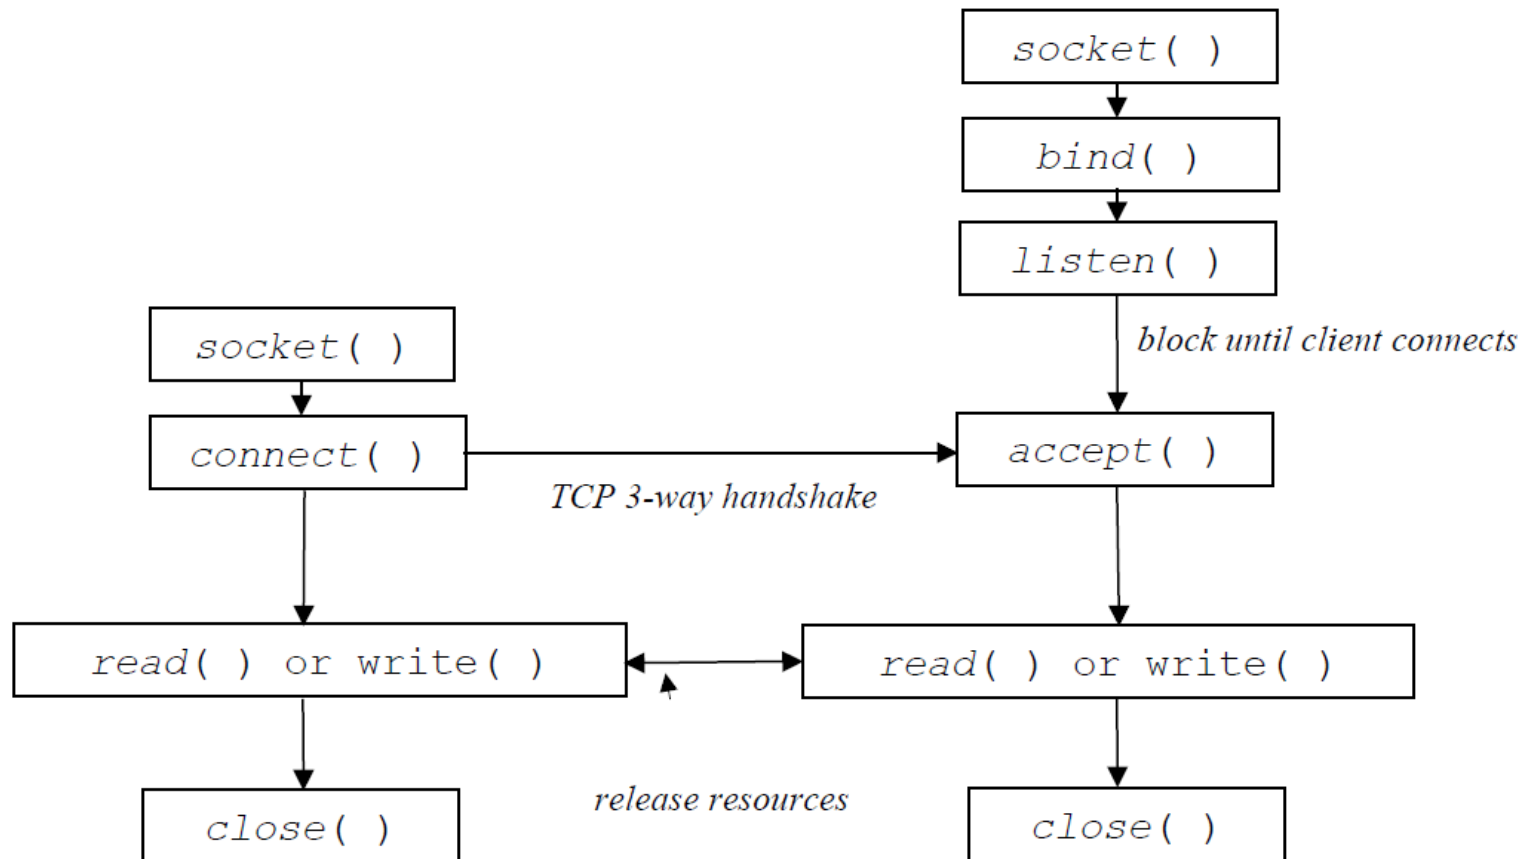
It's Demo Time!

- ▶ Server: Ubuntu, Client: Windows
- ▶ Using *Windows Subsystem for Linux* ([link](#))
 - ▶ For graphical UI under linux (using X11)...
 - ▶ Install (on Windows) Xming X server ([link](#))
 - ▶ Under linux, add to your “.profile” (in home dir):
 - `export DISPLAY=:0`
- ▶ Linux, starting server:
 - ▶ `cd /mnt/c/Artur/OneDrive/AOD/Academic/Lehre/18.WS-IVS1/Code/Lectures/PythonSockets`
 - ▶ `python3 bondServer.py`
- ▶ Windows:
 - ▶ Start from JetBrains's IDEA or Pycharm

TCP Communication in C - Overview

Client

Server



Stream-based Communication in C

Requesting a connection (client)

```
s = socket(AF_INET, SOCK_STREAM, 0)
•
•
connect(s, ServerAddress)
•
•
write(s, "message", length)
```

Listening and accepting a connection (server)

```
s = socket(AF_INET, SOCK_STREAM, 0)
•
bind(s, ServerAddress);
listen(s, 5);
•
sNew = accept(s, ClientAddress);
•
n = read(sNew, buffer, amount)
```

socket()

AF_INET – communication domain is Internet domain

SOCK_STREAM – specifies stream communication

0 – lets the system select the protocol (here TCP)

listen(s, N)

N – maximum number of request for connections to be queued at this socket

Important:
here a new
(„connection“) socket
is created (with its
own port)



Sockets Programming: TCP with Java [Partially Skipped]

Sockets with TCP in Java

▶ Example client-server app:

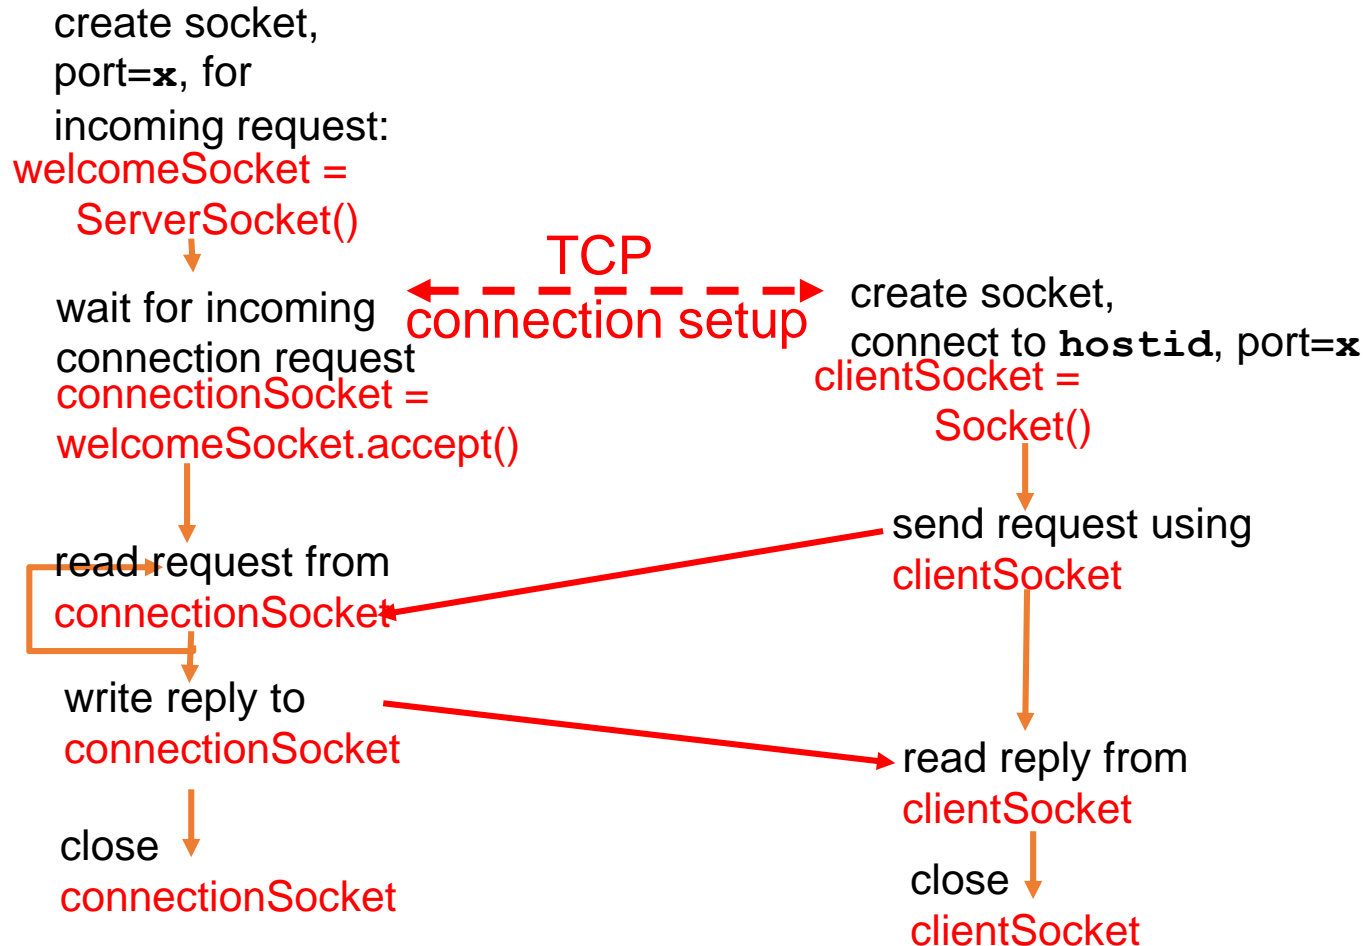
1. client reads line from standard input (**inFromUser** stream), sends to server via socket (**outToServer** stream)
2. server reads line from socket
3. server converts line to uppercase, sends back to client
4. client reads, prints modified line from socket (**inFromServer** stream)



Client/server socket interaction: Java

Server (running on `hostid`)

Client



Example: Java client (TCP)

```
import java.io.*;
import java.net.*;
class TCPCClient {
```

```
    public static void main(String argv[]) throws Exception
    {
```

```
        String sentence;
        String modifiedSentence;
```

Create
input stream



```
        BufferedReader inFromUser =
            new BufferedReader(new InputStreamReader(System.in));
```

Create
client socket,
connect to server



```
        Socket clientSocket = new Socket("hostname", 6789);
```

Create
output stream
attached to socket



```
        DataOutputStream outToServer =
            new DataOutputStream(clientSocket.getOutputStream());
```

Example: Java client (TCP), cont.

Create
input stream
attached to socket

Send line
to server

Read line
from server

```
BufferedReader inFromServer =  
    new BufferedReader(new  
        InputStreamReader(clientSocket.getInputStream()));  
  
sentence = inFromUser.readLine();  
  
outToServer.writeBytes(sentence + '\n');  
  
modifiedSentence = inFromServer.readLine();  
  
System.out.println("FROM SERVER: " + modifiedSentence);  
  
clientSocket.close();  
  
}  
}
```

Example: Java server (TCP)

```
import java.io.*;
import java.net.*;
```

```
class TCPServer {
```

```
    public static void main(String argv[]) throws Exception
    {
```

```
        String clientSentence;
        String capitalizedSentence;
```

Create
welcoming socket
at port 6789

```
        ServerSocket welcomeSocket = new ServerSocket(6789);
```

Wait, on welcoming
socket for contact
by client

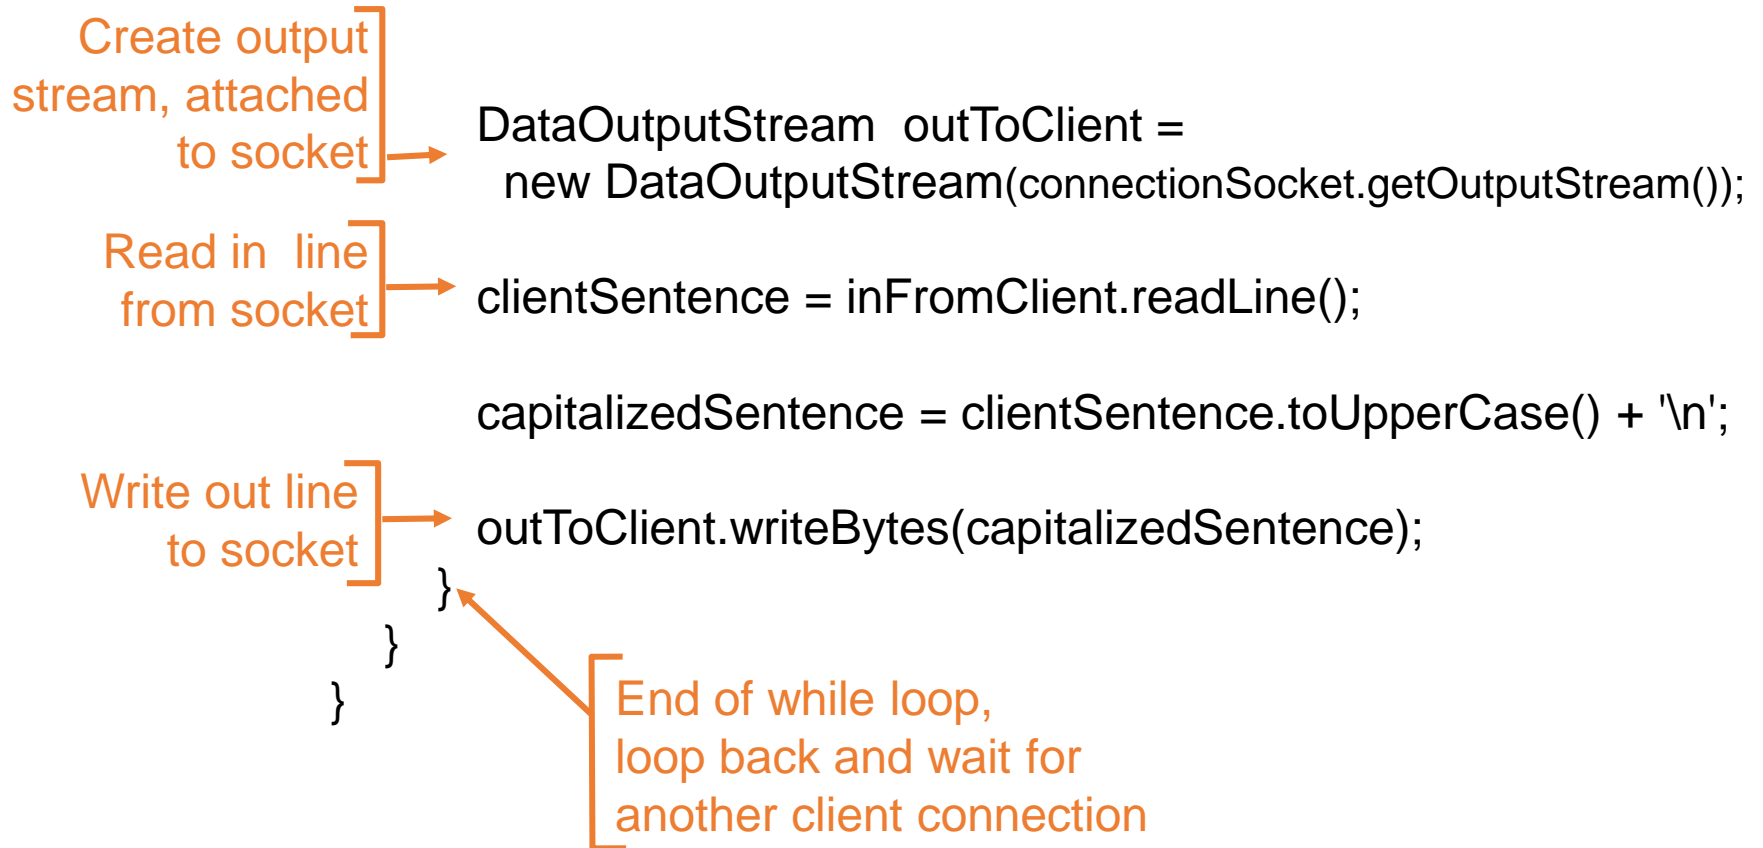
```
        while(true) {
```

```
            Socket connectionSocket = welcomeSocket.accept();
```

Create input
stream, attached
to socket

```
            BufferedReader inFromClient =
                new BufferedReader(new
                    InputStreamReader(connectionSocket.getInputStream()));
```

Example: Java server (TCP), cont



TCP observations & questions

- ▶ Server has two types of sockets:
 - ▶ ServerSocket and Socket
- ▶ When client knocks on serverSocket's "door," server creates connectionSocket and completes TCP connection setup
- ▶ Dest IP and port are not explicitly attached to segment
- ▶ Can multiple clients use the server?

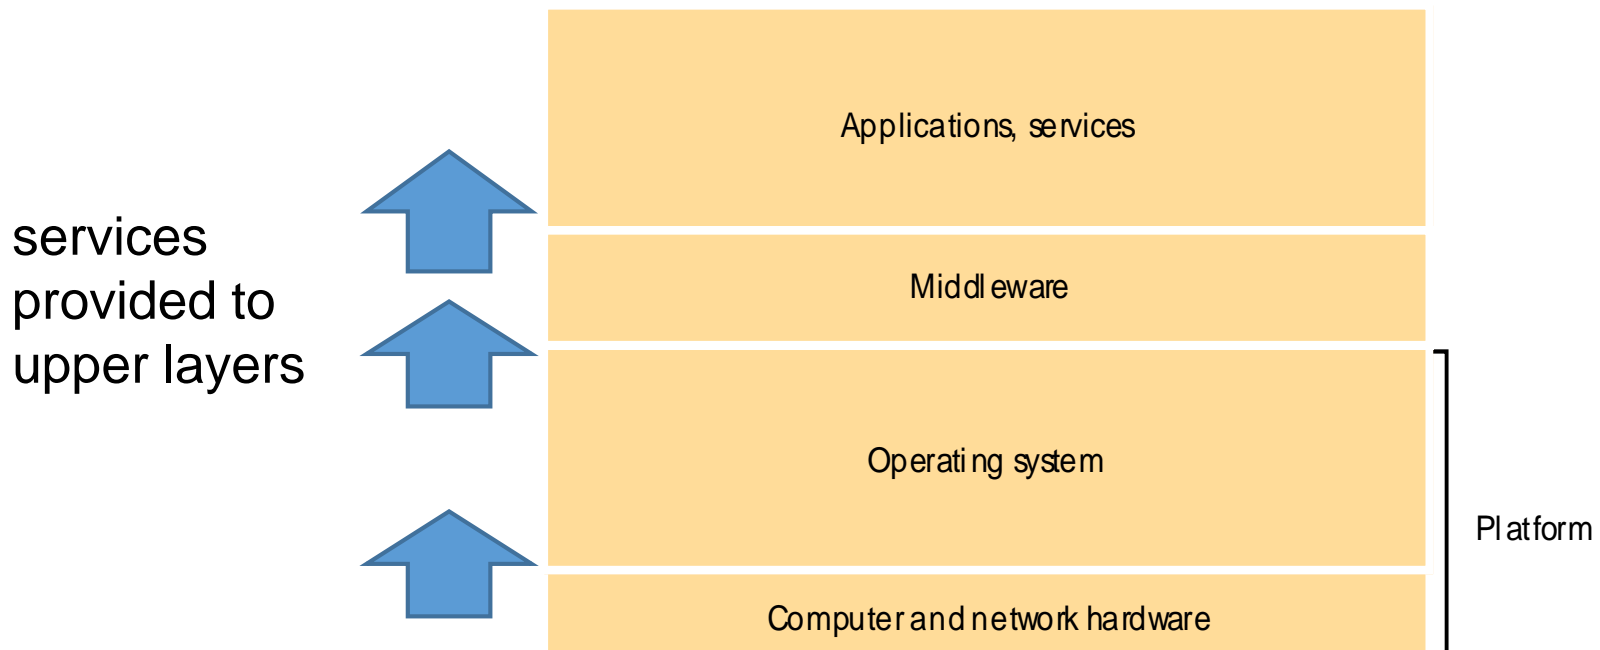


Thank you.

Additional Slides: Architectures and Topologies in Distributed Systems

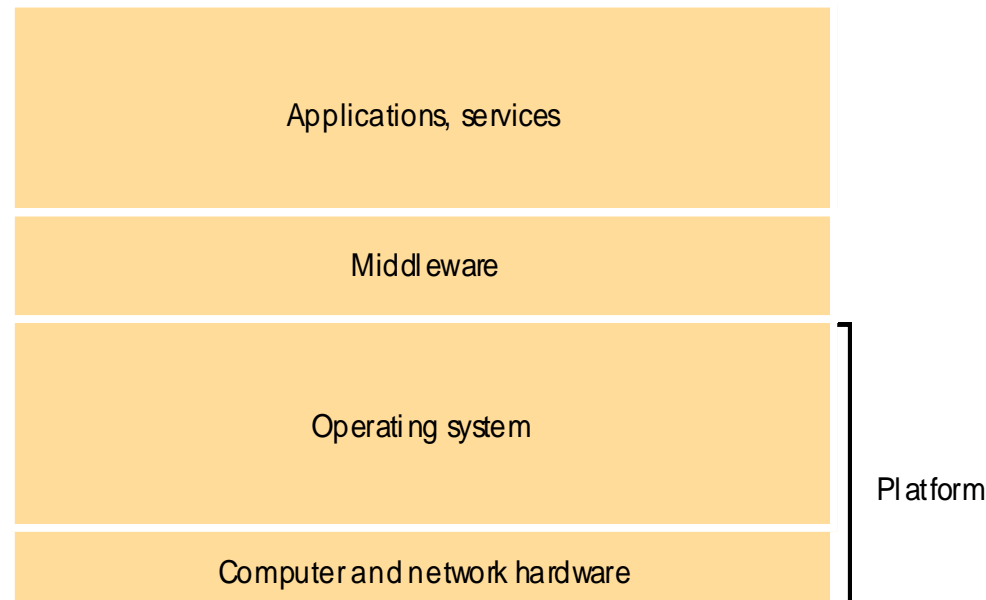
Layers

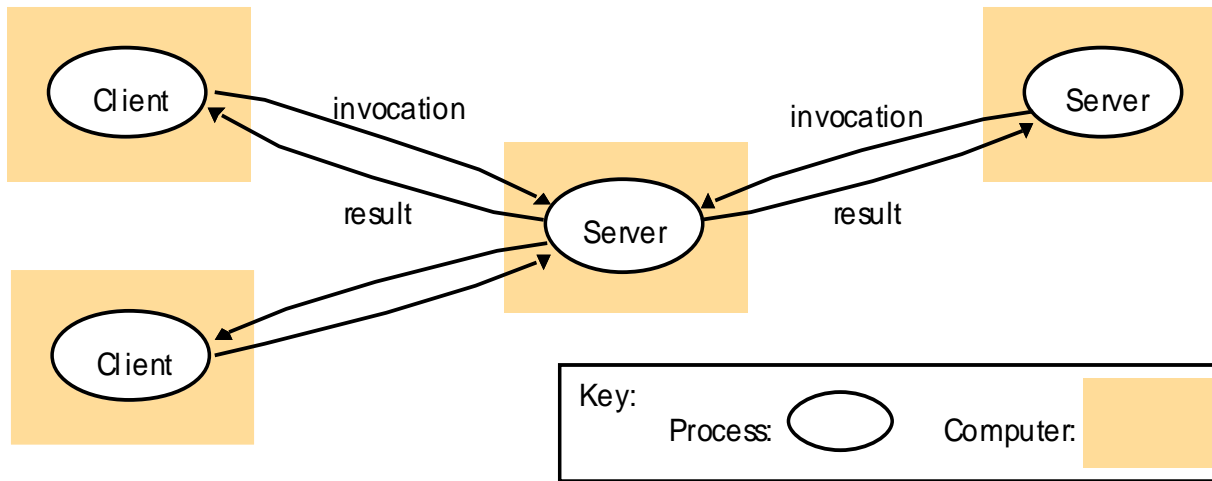
- ▶ Reduce the complexity of systems by designing them through **layers** and **services**
 - ▶ **Layer**: group of closely related and highly coherent functionalities
 - ▶ **Service**: functionality provided to an upper layer



Typical Layers

- ▶ **Platform:** Hardware and operating system
 - ▶ e.g. x86 with Linux vs. SPARC with Solaris
- ▶ **Middleware:** provides communication and resource sharing, e.g.
 - ▶ gRPC / protocol buffers (Google)
 - ▶ Java Remote Method Invocation (Java RMI)



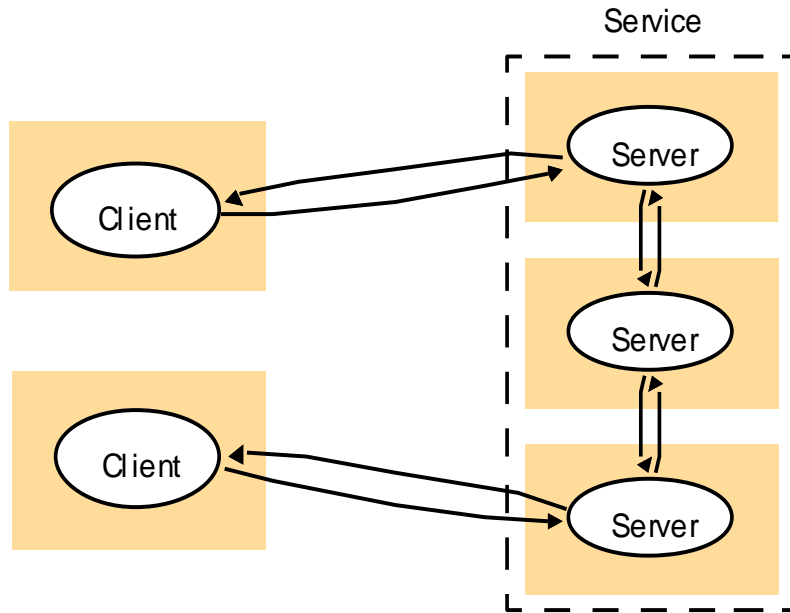


Client – Server Architecture

► Examples:

- http server: client (browser) requests page, server delivers page
- file server: client requests file (or parts of it), server delivers data
- Google apps: storage and processing on Google servers

- **Client**: participant which wants to access data, use resources or perform operations on remote computer
- **Server**: participant managing data and all other shared resources; allows clients access to resource and performs computation



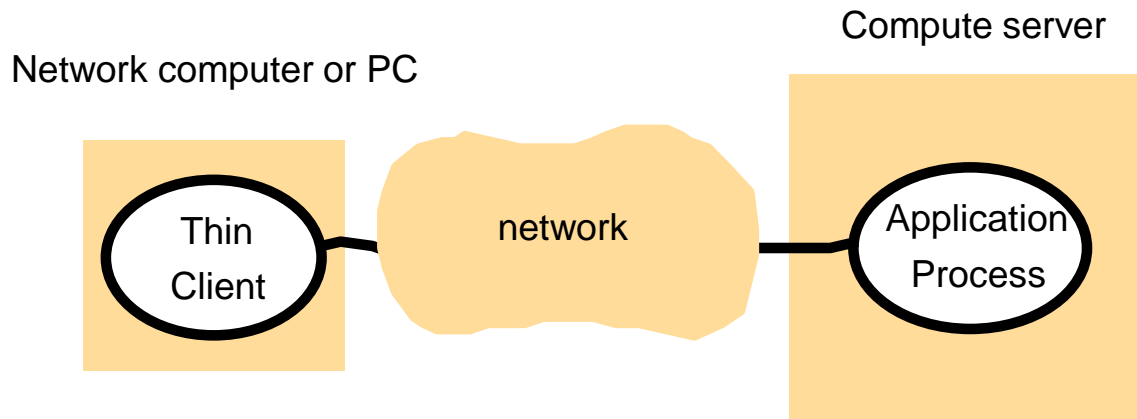
Client – Server Architecture

Variation:

- Service is provided by multiple servers

► Reasons

- Scalability: e.g. Google search provided by thousands of servers
- Fault tolerance: replicated services are more robust against failures



Client – Server Architecture

Thin clients and compute servers

- Thin clients is used merely for presentation and input collecting
- The „ultimate“ form of client/server paradigm

► Reasons

- Low cost of clients
- Continuous software update at server

► Examples

- Google docs / spreadsheets (thin client = browser)
- X11 display protocol in linux/unix
- Dumb terminals in local network (old)

Additional Slides: Interprocess Communication

IPC Types – Example Windows

▶ **Mailslots**

- ▶ Mailslots provide one-way communication
 - ▶ Any process that creates a mailslot is a **mailslot server**
 - ▶ Other processes, called **mailslot clients**, send messages to the mailslot server (by writing a message to its mailslot)
- ▶ The mailslot saves the messages (in FIFO fashion) until the mailslot server has read them
- ▶ A process can be both a mailslot server and a mailslot client, so two-way communication is possible using multiple mailslots



Pipes Example „who | sort“

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>
int main(void){
    int pipe_verbindung[2];
    pipe(pipe_verbindung);
    if (fork()==0){
        dup2(pipe_verbindung[1],1);
        close (pipe_verbindung[0]);
        execlp("who","who",NULL);
    } else if (fork()==0){
        dup2(pipe_verbindung[0],0);
        close(pipe_verbindung[1]);
        execlp("sort","sort",NULL);
    }
}
```

Adapted from Wikipedia:

[http://de.wikipedia.org/wiki/Pipe_\(Informatik\)](http://de.wikipedia.org/wiki/Pipe_(Informatik))

```
// code child process 1
// connect „write-to“ end to std-out
// close the „read-from“ end
// execute „who“

// code child process 2
// connect „read-from“ to std-in
// close the „write-to“ end
```



IPC Types – Mailboxes / MP / RPC

- ▶ **Message queue / mailbox:** most OSs
 - ▶ Typically performed by an off-the-shelf message-queuing software
 - ▶ application registers a routine that "listens" for messages placed onto the queue
 - ▶ second and subsequent applications may connect to the queue and transfer a message onto it
- ▶ **Message passing and Remote Procedure Calls:**
 - ▶ Via Middleware (MPI, Java RMI, CORBA, ...)
 - ▶ Communication via sending of messages to recipients
 - ▶ (Much) more on this later



IPC Types – Example Windows

- ▶ **Clipboard**
- ▶ **COM** and **OLE** – (COM = Component Object Model)
 - ▶ “OLE supports compound documents and enables an application to include embedded or linked data that, when chosen, automatically starts another application for data editing.”
 - ▶ “COM objects provide access to an object's data through one or more sets of related functions, known as interfaces.”
- ▶ **Data Copy**
 - ▶ „Data copy enables an application to send information to another application using the WM_COPYDATA message.”
 - ▶ “This method requires cooperation between the sending application and the receiving application.”



IPC Types – Example Windows

- ▶ **DDE** (Dynamic Data Exchange)
 - ▶ DDE is an extension of the clipboard mechanism
 - ▶ It is also usually initiated by a user command, but it often continues to function without further user interaction
 - ▶ DDE is not as efficient as newer technologies
- ▶ **File Mapping**
 - ▶ See Memory-mapped file (mmap)
- ▶ **Anonymous Pipes**
- ▶ **Named Pipes**



Additional Slides: UDP Sockets in Java

Running Example

▶ Client:

- ▶ User types line of text
- ▶ Client program sends line to server

▶ Server:

- ▶ Server receives line of text
- ▶ Capitalizes all the letters
- ▶ Sends modified line to client

▶ Client:

- ▶ Receives line of text
- ▶ Displays



Client/server Socket Interaction: UDP

Server (running on `hostid`)

Client

create socket,
port= x.
`serverSocket =`
`DatagramSocket(x)`

read datagram from
`serverSocket`

write reply to
`serverSocket`
specifying
client address,
port number

create socket,
`clientSocket =`
`DatagramSocket()`

Create datagram with server IP and
port=x; send datagram via
`clientSocket`

read datagram from
`clientSocket`

close
`clientSocket`

Example: Java client (UDP)

```
import java.io.*;
import java.net.*;
```

```
class UDPClient {
    public static void main(String args[]) throws Exception
    {
```

Create
input stream

```
        BufferedReader inFromUser =
            new BufferedReader(new InputStreamReader(System.in));
```

Create
client socket

```
        DatagramSocket clientSocket = new DatagramSocket();
```

Translate
hostname to IP
address using DNS

```
        InetAddress IPAddress = InetAddress.getByName("hostname");
```

```
        byte[] sendData = new byte[1024];
        byte[] receiveData = new byte[1024];
```

```
        String sentence = inFromUser.readLine();
        sendData = sentence.getBytes();
```

Example: Java client (UDP), cont.

Create datagram with
data-to-send,
length, IP addr, port

```
DatagramPacket sendPacket =  
    new DatagramPacket(sendData, sendData.length,  
                        IPAddress, 9876);
```

Send datagram
to server

```
clientSocket.send(sendPacket);
```

Read datagram
from server

```
DatagramPacket receivePacket =  
    new DatagramPacket(receiveData, receiveData.length);
```

```
clientSocket.receive(receivePacket);
```

```
String modifiedSentence =  
    new String(receivePacket.getData());
```

```
System.out.println("FROM SERVER:" + modifiedSentence);  
clientSocket.close();
```

```
}
```

```
}
```


Example: Java server (UDP)

```
import java.io.*;  
import java.net.*;
```

```
class UDPServer {  
    public static void main(String args[]) throws Exception  
    {
```

Create
datagram socket
at port 9876

```
        DatagramSocket serverSocket = new DatagramSocket(9876);
```

```
        byte[] receiveData = new byte[1024];  
        byte[] sendData = new byte[1024];
```

```
        while(true)  
        {
```

Create space for
received datagram

```
            DatagramPacket receivePacket =  
                new DatagramPacket(receiveData, receiveData.length);
```

Receive
datagram

```
            serverSocket.receive(receivePacket);
```

Example: Java server (UDP), cont

```
String sentence = new String(receivePacket.getData());
```

Get IP addr
port #, of
sender

```
InetAddress IPAddress = receivePacket.getAddress();  
int port = receivePacket.getPort();
```

```
String capitalizedSentence = sentence.toUpperCase();
```

```
sendData = capitalizedSentence.getBytes();
```

Create datagram
to send to client

```
DatagramPacket sendPacket =  
    new DatagramPacket(sendData, sendData.length, IPAddress,  
                        port);
```

Write out
datagram
to socket

```
serverSocket.send(sendPacket);
```

```
}  
}  
}
```

End of while loop,
loop back and wait for
another datagram

UDP Observations & Questions

- ▶ Both client server use DatagramSocket
- ▶ Dest IP and port are explicitly attached to segment
- ▶ Can the client send a segment to server without knowing the server's IP address and/or port number?
- ▶ Can multiple clients use the server?

