

Verteilte Systeme/ Distributed Systems

Artur Andrzejak

7

Time and Clock Synchronization

Following contents are based on the slides of Lecture 09 in the course:
Distributed Computing, Google Code University/ Rutgers University:
By Paul Krzyzanowski, pxk@cs.rutgers.edu, ds@pk.org
Attribution according to Creative Commons Attribution 2.5 License.

What's it for?

- ▶ Temporal ordering of events produced by concurrent processes
- ▶ Synchronization between senders and receivers of messages
- ▶ Coordination of joint activity
- ▶ Serialization of concurrent access for shared objects

Physical Clocks in Computers

- ▶ Real-time Clock: CMOS clock (counter) circuit driven by a quartz oscillator
 - ▶ Battery backup to continue measuring time when power is off
- ▶ OS generally programs a timer circuit to generate an interrupt periodically
 - ▶ e.g., 60, 100, 250, 1000 interrupts per second (Linux 2.6+ adjustable up to 1000 Hz)
 - ▶ Programmable Interval Timer (PIT) – Intel 8253, 8254
 - ▶ Interrupt service procedure adds 1 to a counter in memory

Problems

- ▶ Getting two systems to agree on time
 - ▶ Two clocks hardly ever agree
 - ▶ Quartz oscillators oscillate at slightly different frequencies
- ▶ Clocks tick at different rates
 - ▶ Create ever-widening gap in perceived time
 - ▶ **Clock Drift**
- ▶ Difference between two clocks at one point in time
 - ▶ **Clock Skew**

Dealing with Drift

Following contents are based on the slides of Lecture 09 in the course:
Distributed Computing, Google Code University/ Rutgers University:
By Paul Krzyzanowski, pxk@cs.rutgers.edu, ds@pk.org
Attribution according to Creative Commons Attribution 2.5 License.



8:00:00



8:00:00

Sept 18, 2006
8:00:00



8:01:24

Skew = +84 seconds
 +84 seconds/35 days
 Drift = +2.4 sec/day




8:01:48

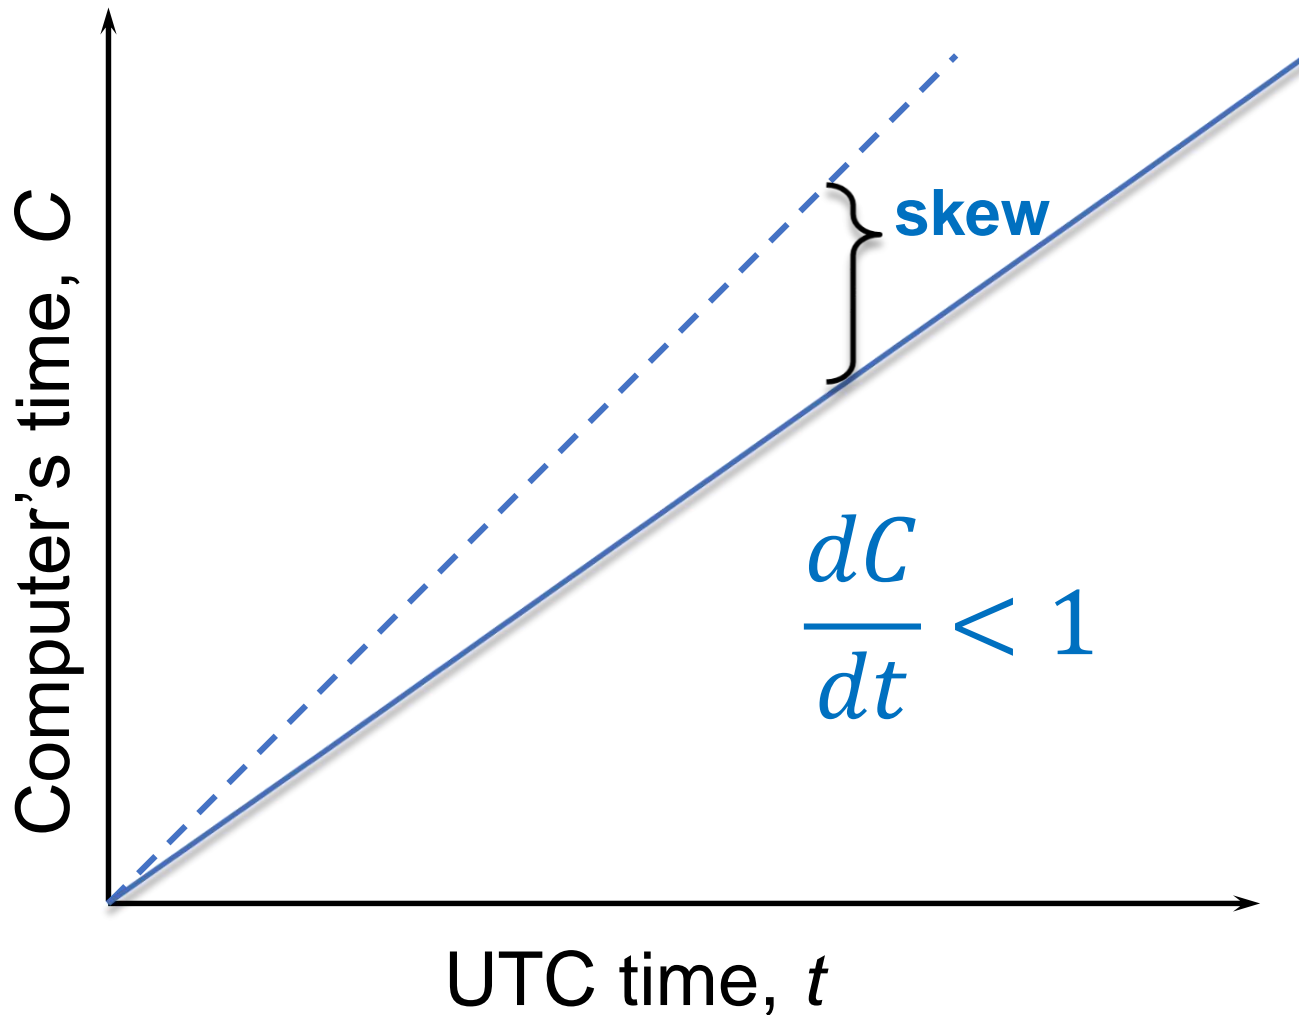
Skew = +108 seconds
 +108 seconds/35 days
 Drift = +3.1 sec/day

Oct 23, 2006
 8:00:00

Perfect Clock



Drift with Slow Clock



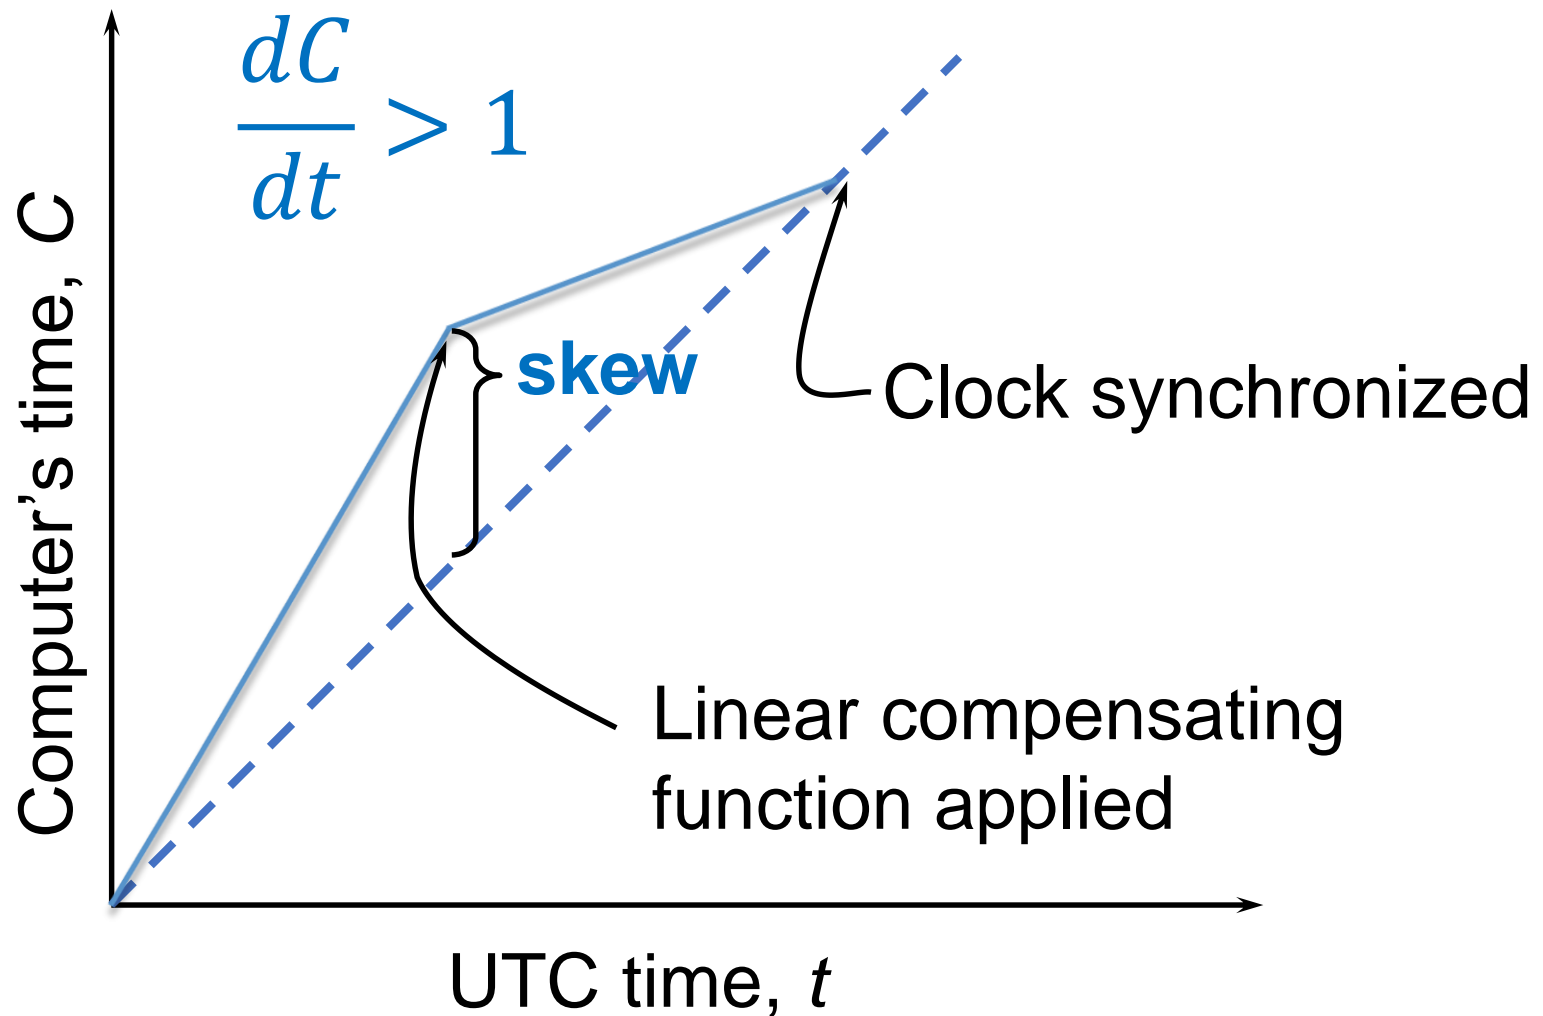
Dealing With Drift

- ▶ Assume we set computer to true time
- ▶ Not good idea to set clock back
 - ▶ Illusion of time moving backwards can confuse message ordering and software development environments!
- ▶ How to do it better?
- ▶ Go for gradual clock correction
 - ▶ If fast: make clock run slower until it synchronizes
 - ▶ If slow: make clock run faster until it synchronizes

Dealing With Drift

- ▶ OS can do this:
 - ▶ Change rate at which it requests interrupts, e.g.:
 - ▶ if system requests interrupts every 17 msec but clock is too slow:
 - ▶ request interrupts at (e.g.) 15 msec
 - ▶ Or software correction: redefine the interval
- ▶ Adjustment changes slope of system time:
 - ▶ Linear compensating function

Compensating for a Fast Clock



Resynchronizing

- ▶ After synchronization period is reached
 - ▶ Resynchronize periodically
 - ▶ Successive application of a second linear compensating function can bring us closer to true slope
- ▶ Keep track of adjustments and apply continuously
 - ▶ e.g., UNIX *adjtime* system call

Getting Accurate Time

Following contents are based on the slides of Lecture 09 in the course:
Distributed Computing, Google Code University/ Rutgers University:
By Paul Krzyzanowski, pxk@cs.rutgers.edu, ds@pk.org
Attribution according to Creative Commons Attribution 2.5 License.


Getting Accurate time

- ▶ Attach GPS receiver to each computer
 - ▶ ± 1 msec of UTC
- ▶ Attach WWV radio receiver ([link](#))
 - ▶ Obtain time broadcasts from Boulder or DC
 - ▶ ± 3 msec of UTC (depending on distance)
- ▶ Attach GOES receiver ([link](#))
 - ▶ ± 0.1 msec of UTC
- ▶ Not practical solution for every machine – what else?
- ▶ Synchronize from another machine
 - ▶ One with a more accurate clock
- ▶ Machine/service that provides time information:
 - ▶ **Time server**

RPC

Simplest synchronization technique


- ▶ Issue RPC to obtain time
- ▶ Set time



- ▶ Problems?
- ▶ Does not account for network or processing latency

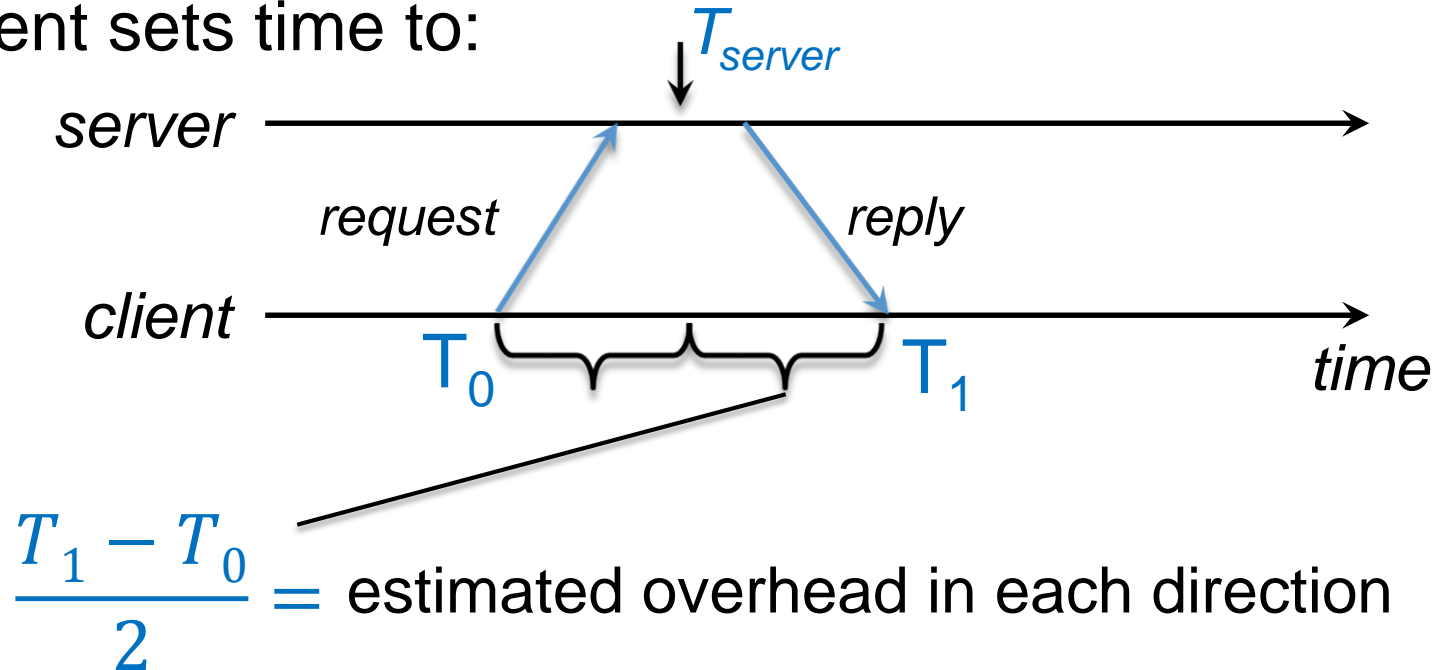
Cristian's algorithm

- ▶ Compensate for delays
 - ▶ Note times:
 - ▶ request sent: T_0
 - ▶ reply received: T_1
 - ▶ Assume network delays are symmetric



Cristian's Algorithm


- ▶ Client sets time to:



$$T_{new} = T_{server} + \frac{T_1 - T_0}{2}$$

Cristian's Algorithm - Error Bounds:

If minimum message transit time (T_{\min}) is known



accuracy of result = $\pm \frac{T_1 - T_0}{2} - T_{\min}$

Berkeley Algorithm

- ▶ Gusella & Zatti, 1989
- ▶ Assumes no machine has an accurate time source
- ▶ Obtains average from participating computers
- ▶ Synchronizes all clocks to average

- ▶ Machines run **time daemon**
 - ▶ Process that implements protocol
- ▶ One machine is elected (or designated) as the server (**master**)
 - ▶ Others are **slaves**

Berkeley Algorithm


- ▶ Master polls each machine periodically, i.e. ask each machine for time
 - ▶ Can use Cristian's algorithm to compensate for network latency
- ▶ When results are in, compute average, including master's
- ▶ Hope: average cancels out individual clock's tendencies to run fast or slow
- ▶ Improvements:
 - ▶ Send offset by which each clock needs adjustment to each slave
 - ▶ Avoids problems with network delays if we send a time stamp
 - ▶ Algorithm has provisions for ignoring readings from clocks whose skew is too great
 - ▶ Compute a fault-tolerant average
 - ▶ If master fails: Any slave can take over

Network Time Protocol, **NTP**

- ▶ 1991, 1992: Internet Standard, version 3: RFC 1305
- ▶ Enable clients across Internet to be accurately synchronized to UTC despite message delays
 - ▶ Use statistical techniques to filter data and gauge quality of results
- ▶ Provide reliable service
 - ▶ Survive lengthy losses of connectivity
 - ▶ Redundant paths
 - ▶ Redundant servers
- ▶ Enable clients to synchronize frequently
 - ▶ Offset effects of clock drift
- ▶ Provide protection against interference
 - ▶ Authenticate source of data

NTP Servers

- ▶ ... Arranged in **strata**
 - ▶ 1st stratum: machines connected directly to accurate time source
 - ▶ 2nd stratum: machines synchronized from 1st stratum machines
 - ▶ ...



SYNCHRONIZATION
SUBNET

NTP Synchronization Modes

- ▶ Multicast mode
 - ▶ for high speed LANS
 - ▶ Lower accuracy but efficient
- ▶ Procedure call mode
 - ▶ Similar to Cristian's algorithm
- ▶ Symmetric mode
 - ▶ Intended for master servers
 - ▶ Pair of servers exchange messages and retain data to improve synchronization over time

All messages delivered unreliably with UDP

NTP Messages


- ▶ Procedure call and symmetric mode
 - ▶ Messages exchanged in pairs
- ▶ NTP calculates:
 - ▶ **Offset** for each pair of messages
 - ▶ Estimate of offset between two clocks
 - ▶ **Delay**
 - ▶ Transmit time between two messages
 - ▶ **Filter Dispersion**
 - ▶ Estimate of error – quality of results
 - ▶ Based on accuracy of server's clock and consistency of network transit time
- ▶ Use this data to find preferred server:
 - ▶ Lower stratum & lowest total dispersion

SNTP

Simple Network Time Protocol

- ▶ Based on Unicast mode of NTP
 - ▶ Subset of NTP, not new protocol
 - ▶ Operates in multicast or procedure call mode
 - ▶ Recommended for environments where server is root node and client is leaf of synchronization subnet
 - ▶ Root delay, root dispersion, reference timestamp ignored
-
- ▶ RFC 2030, October 1996

SNTP




Roundtrip delay:

$$d = (T_4 - T_1) - (T_2 - T_3)$$

Time offset:

$$t = \frac{(T_2 - T_1) - (T_3 - T_4)}{2}$$

SNTP Example



Offset =

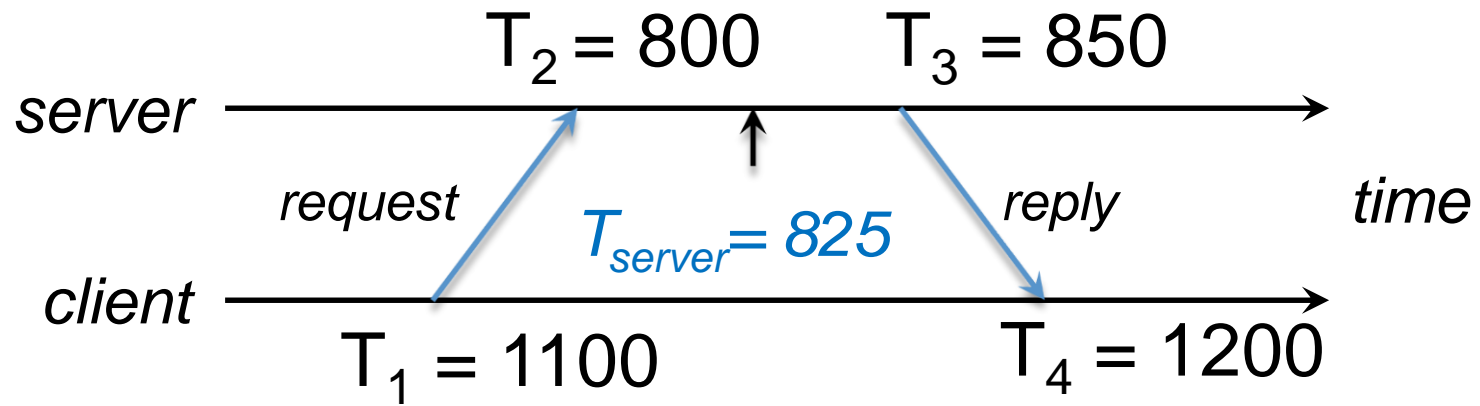
$$\begin{aligned} & ((800 - 1100) + (850 - 1200))/2 \\ & = ((-300) + (-350))/2 \\ & = -650/2 = -325 \end{aligned}$$

Time offset:

$$t = \frac{(T_2 - T_1) - (T_3 - T_4)}{2}$$

Set time to $T_4 + t$
 $= 1200 - 325 = 875$

Compare to Christian's Algorithm



Offset =

$$(1200 - 1100)/2 = 50$$

Time offset:

$$t = \frac{(T_2 - T_1) - (T_3 - T_4)}{2}$$

Set time to $T_{\text{server}} + \text{offset}$
 $= 825 + 50 = 875$

Key Points: Physical Clocks

- ▶ Cristian's algorithm & SNTP
 - ▶ Set clock from server
 - ▶ But account for network delays
 - ▶ Error: uncertainty due to network/processor latency: errors are additive
 $\pm 10 \text{ msec}$ and $\pm 20 \text{ msec} = \pm 30 \text{ msec}$.
- ▶ Adjust for local clock skew
 - ▶ Linear compensating function

Logical Time: Introduction


Logical Time

- ▶ Physical clocks can not be perfectly synchronized
- ▶ What we often only need is a unique sequence of events in a distributed system, not absolute time
 - ▶ But based on physical clocks we can not define a clear order
- ▶ Solution by Leslie Lamport: **logical time** (definition follows)
 - ▶ Paper: "Time, Clocks, and the Ordering of Events in a Distributed System," Commun. ACM, vol. 21, no. 7, pp. 558-565, 1978
- ▶ We first introduce the **temporal order** through two rules:
 - ▶ If two events take place in the same process P, then they have the order (order) that P has observed
 - ▶ When a message is sent, the event of transmission occurs before the event of reception

Happened-Before Relation

- ▶ Lamport defined from the two relationships the **Happened-Before** relation \rightarrow
 - ▶ It formalizes the temporal order
- ▶ Happened-Before-Relation \rightarrow is defined by:
 - ▶ **HB1**: If there is a process p_i with $e \rightarrow_i e'$, then $e \rightarrow e'$
 - ▶ I.e. the "local" order \rightarrow_i in p_i specifies \rightarrow for e, e'
 - ▶ **HB2**: For each message m : $\text{send}(m) \rightarrow \text{receive}(m)$
 - ▶ I.e. event $\text{send}(m)$ is before event $\text{receive}(m)$
 - ▶ **HB3**: If for e, e', e'' we have: $e \rightarrow e'$ and $e' \rightarrow e''$, then $e \rightarrow e''$ (transitivity)

Examples and Concurrent Events



- ▶ Example: $a \rightarrow b$, $b \rightarrow c$, $a \rightarrow f$
- ▶ For some events the HB relation is not true - examples?
- ▶ Example: $a \nrightarrow e$, $e \nrightarrow a$
- ▶ Events x , y that are not ordered with \rightarrow are called **concurrent**, denoted as $x \parallel y$


Lamport Clock

- ▶ Lamport (1978) developed a simple method to numerically express Happened-Before-Relation: **Lamport clock** (logical clock)
- ▶ Lamport clock is a monotonously increasing (software) counter
 - ▶ No relationship to the physical time needed
- ▶ Each process p_i has its own logical clock L_i
- ▶ Process p_i assigns timestamps to events by:
 - ▶ Timestamp of the clock L_i (L_i local for p_i) for event e is the Lamport timestamp **$L_i(e)$** (note: index i in L_i)
- ▶ **$L(e)$** (without index i) is the Lamport timestamp of an event, regardless of where it occurred

Lamport Clock and Happened-Before

- ▶ To comply with the Happened-Before relationship, processes change their logical clock as follows:
- ▶ **LC1:**
 - ▶ L_i is incremented just before an event occurs on p_i , i.e. $L_i := L_i + 1$
- ▶ **LC2:**
 - ▶ a) When a process p_i sends a message m , it also sends with m the actual value **t** of its own logical clock L_i
 - ▶ b) When other process p_j receives the values (m, \mathbf{t})
 - ▶ b1. Calculate $\max(L_j, \mathbf{t})$, and set its new L_j to this value, i.e. $L_j := \max(L_j, \mathbf{t})$
 - ▶ b2. Executes LC1 (i.e., increments L_j)
 - ▶ b3. Then assign to $\text{receive}(m)$ the timestamp computed in b2


Lamport Clock: Quiz and a Surprise



- ▶ Each process has its own logical clock, initialized with 0
- ▶ What are $L(a)$, $L(b)$, ..., $L(f)$?
- ▶ Note the following: $L(b) > L(e)$, but $b \parallel e$

Lamport Clock: Problems

- ▶ Lamport clock ensures that ...
 - ▶ For two events a and b with $a \rightarrow b$ (a happened-before b) we have: $L(a) < L(b)$
- ▶ But the converse does not apply!
 - ▶ From $L(a) < L(b)$ it does not follow that $a \rightarrow b$ holds
- ▶ Therefore, in general we cannot conclude about the causality of two events based on the Lamport time only



We have $L(d) < L(b)$,
but not $d \rightarrow b$
(it holds: $d \parallel b$)

Thank you.

Additional Slides

Additional Slides: Time and Clock Synchronization

Following contents are based on the slides of Lecture 09 in the course:
Distributed Computing, Google Code University/ Rutgers University:
By Paul Krzyzanowski, pxk@cs.rutgers.edu, ds@pk.org
Attribution according to Creative Commons Attribution 2.5 License.

What's it for?

- ▶ Temporal ordering of events produced by concurrent processes
- ▶ Synchronization between senders and receivers of messages
- ▶ Coordination of joint activity
- ▶ Serialization of concurrent access for shared objects

Quartz clocks

▶ 1880: Piezoelectric effect

- ▶ Curie brothers
- ▶ Squeeze a quartz crystal & it generates an electric field
- ▶ Apply an electric field and it bends

▶ 1929: Quartz crystal clock

- ▶ Resonator shaped like tuning fork
- ▶ Laser-trimmed to vibrate at 32,768 Hz
- ▶ Standard resonators accurate to 6 parts per million at 31° C
- ▶ Watch will gain/lose $< \frac{1}{2}$ sec/day
- ▶ Stability > accuracy: stable to 2 sec/month
- ▶ Good resonator can have accuracy of 1 second in 10 years
 - ▶ Frequency changes with age, temperature, and acceleration

Atomic clocks

- ▶ Second is defined as 9,192,631,770 periods of radiation corresponding to the transition between two hyperfine levels of cesium-133
- ▶ Accuracy:
better than 1 second in six million years
- ▶ NIST standard since 1960

UTC

- ▶ UT0
 - ▶ Mean solar time on Greenwich meridian
 - ▶ Obtained from astronomical observation
- ▶ UT1
 - ▶ UT0 corrected for polar motion
- ▶ UT2
 - ▶ UT1 corrected for seasonal variations in Earth's rotation
- ▶ UTC
 - ▶ Civil time measured on an atomic time scale

UTC

- ▶ Coordinated Universal Time
- ▶ Temps Universel Coordonné
 - ▶ Kept within 0.9 seconds of UT1
 - ▶ Atomic clocks cannot keep mean time
 - ▶ Mean time is a measure of Earth's rotation


Cristian's algorithm: example

- ▶ Send request at 5:08:15.100 (T_0)
- ▶ Receive response at 5:08:15.900 (T_1)
 - ▶ Response contains 5:09:25.300 (T_{server})
- ▶ Elapsed time is $T_1 - T_0$
5:08:15.900 - 5:08:15.100 = 800 msec
- ▶ Best guess: timestamp was generated
400 msec ago
- ▶ Set time to $T_{server} + \text{elapsed time}$
5:09:25.300 + 400 = 5:09:25.700

Cristian's algorithm: example


If best-case message time=200 msec

$T_0 = 5:08:15.100$
 $T_1 = 5:08:15.900$
 $T_s = 5:09:25:300$
 $T_{min} = 200\text{msec}$




$$\text{Error} = \pm \frac{900 - 100}{2} - 200 = \pm \frac{800}{2} - 200 = \pm 200$$

Berkeley Algorithm: example



1. Request timestamps from all slaves


Berkeley Algorithm: example



2. Compute fault-tolerant average:

$$\frac{3:25 + 2:50 + 3:00}{3} = 3:05$$

Berkeley Algorithm: example



3. Send offset to each client

NTP Message Structure

- ▶ Leap second indicator
 - ▶ Last minute has 59, 60, 61 seconds
- ▶ Version number
- ▶ Mode (symmetric, unicast, broadcast)
- ▶ Stratum (1=primary reference, 2-15)
- ▶ Poll interval
 - ▶ Maximum interval between 2 successive messages, nearest power of 2
- ▶ Precision of local clock
 - ▶ Nearest power of 2

NTP Message Structure

- ▶ Root delay
 - ▶ Total roundtrip delay to primary source
 - ▶ (16 bits seconds, 16 bits decimal)
- ▶ Root dispersion
 - ▶ Nominal error relative to primary source
- ▶ Reference clock ID
 - ▶ Atomic, NIST dial-up, radio, LORAN-C navigation system, GOES, GPS, ...
- ▶ Reference timestamp
 - ▶ Time at which clock was last set (64 bit)
- ▶ Authenticator (key ID, digest)
 - ▶ Signature (ignored in SNTP)

NTP Message Structure

- ▶ T_1 : originate timestamp
 - ▶ Time request departed client (client's time)
- ▶ T_2 : receive timestamp
 - ▶ Time request arrived at server (server's time)
- ▶ T_3 : transmit timestamp
 - ▶ Time request left server (server's time)

NTP's Validation Tests

- ▶ Timestamp provided \neq last timestamp received
 - ▶ duplicate message?
- ▶ Originating timestamp in message consistent with sent data
 - ▶ Messages arriving in order?
- ▶ Timestamp within range?
- ▶ Originating and received timestamps $\neq 0$?
- ▶ Authentication disabled? Else authenticate
- ▶ Peer clock is synchronized?
- ▶ Don't sync with clock of higher stratum #
- ▶ Reasonable data for delay & dispersion

Logical Time: Backup Slides

Following contents are based on the slides and book:
Ajay Kshemkalyani and Mukesh Singhal,
Distributed Computing: Principles, Algorithms, and Systems,
Cambridge University Press, 2011, [link](#)

Causality and Time

- ▶ The concept of **causality** between events is fundamental to the design and analysis of parallel and distributed computing and operating systems
- ▶ Usually causality is tracked using **physical time**
 - ▶ In distributed systems, it is not possible to have a (correct) global physical time
- ▶ Asynchronous distributed computations make progress in spurts ...
- ▶ .. So the so-called **logical time** is sufficient to capture the fundamental monotonicity property associated with causality in distributed systems