

Verteilte Systeme/ Distributed Systems

Artur Andrzejak

10

Self-stabilizing Systems

Stabilizing Algorithms - Intuition

- ▶ An algorithm is **stabilizing** if it can finally bring the system to a correct state (according to the specification) regardless of the initial configuration
- ▶ Stabilizing algorithms are "optimistic"
 - ▶ They assume that intermittent errors can occur, which then disappear (**transient errors**)
 - ▶ The correct processes react without checks to the (incorrect) messages, which can cause arbitrary corrupt system states
 - ▶ However, it is assumed that any process will eventually return to normal execution
 - ▶ The "stabilization property" of the algorithm causes the system to finally transition to a correct configuration (= state)




System, Execution, Specification

- ▶ A **system** is a pair $S=(C, \rightarrow)$, where
 - ▶ C is a set of **configurations** (i.e. global states)
 - ▶ “ \rightarrow ” is a binary relation on C , a **state transition relation**
 - ▶ I.e. $g_i \rightarrow g_{i+1}$ holds if one can move from the state g_i to g_{i+1}
- ▶ An **execution** of S is a maximal sequence g_0, g_1, \dots , so that:
 - ▶ $g_i \rightarrow g_{i+1}$ holds for all $i \geq 0$
- ▶ The desired consistent behavior of the processes is modeled as a **predicate P** on the sequences of configurations
- ▶ **P** is called a **specification**
 - ▶ Example of P : "each configuration is the same"
 - ▶ I.e. one execution satisfies P when system enters a configuration and never leaves it again
 - ▶ Note: to verify P , we must consider a sequence of configurations!

Stabilization

- ▶ A system S **stabilizes** to the specification P if there is a set $L \subseteq C$ of so-called **legitimate configurations** with the following properties:
- ▶ Correctness:
 - ▶ Any execution starting from a configuration in L satisfies P
 - ▶ I.e. Each configuration in L can be seen as a “reset to a clean state”, after which P holds
- ▶ Convergence:
 - ▶ Each execution contains a configuration of L
 - ▶ I.e. No matter how we execute, we will “pass” through a “cleaning configuration” from L

Example: a Stabilizing System



1. Correctness?
2. Convergence?

Legitimate configurations

Every execution over these configurations fulfills **P**

Good Properties of Stabilizing Systems

▶ Fault tolerance

- ▶ Full and automatic protection against all transient process errors, because the system finds itself in a correct state from any - and also so corrupt - configuration

▶ Initialization

- ▶ The need for consistent initialization is eliminated because the processes can start in arbitrary states and ultimately achieve coordinated behavior

▶ Dynamic topology

- ▶ A stabilizing algorithm that is topology-dependent (e.g. uses routing tables) converges to a new solution after the occurrence of a topological change

Bad Properties of Stabilizing Systems

- ▶ **Initial inconsistency**

- ▶ Before a *legitimate configuration* is achieved, the algorithm may have an inconsistent output

- ▶ **Inefficiency**


- ▶ Self-stabilizing algorithms are typically less efficient than classical algorithms for the same problem

- ▶ **Ignorance of the stable state**

- ▶ It is not possible to determine within the system that a legitimate configuration has been achieved
- ▶ Therefore, the processes can not tell whether their behavior has become reliable

Pseudo-Stabilizing Systems

- ▶ Theorem : *If a system stabilizes to **P**, then each execution has a non-empty suffix that satisfies **P**.*
 - ▶ Suffix = final part („Nachsilbe“, „Endsilbe“)
- ▶ A system is **pseudo-stabilizing** to **P** if each execution has a non-empty suffix that satisfies **P**
- ▶ Example
 - ▶ System S with configs a, b and transitions $a \rightarrow a$, $a \rightarrow b$, $b \rightarrow b$
 - ▶ Let **P** be the predicate "*all configurations are the same*"
 - ▶ Is S pseudo-stabilizing? Is S stabilizing?



Each execution has a suffix of the same configuration => S is pseudo-stabilizing 9

Pseudo-Stabilizing Systems

- ▶ **Pseudo-stabilization**: A weaker property than stabilizing systems
- ▶ Good: one can find pseudo-stabilizing algorithms for problems that do not have stabilizing algorithms
 - ▶ Example: Transferring a sequence of data
- ▶ Bad: For some pseudo-stabilizing systems there is no limit on the number of steps to reach the correct state
- ▶ For stabilizing systems such a barrier can always be specified

Mutual Exclusion / Wechselseitiger Ausschluss

Recall: Dijkstra's Token Ring

System model and conventions

- ▶ Let N be the number of processes
- ▶ A state s_i of process i is an integer $s_i \in \{0, \dots, K - 1\}$
 - ▶ Here K is an arbitrary integer larger N ($K > N$)
- ▶ Process $i > 0$ can read s_{i-1} (and of course s_i)
- ▶ Process $i = 0$ can read s_{N-1} (and of course s_0)

Privilege rules (i.e. right to execute critical section):

- ▶ Process $i > 0$ has a privilege iff $s_{i-1} \neq s_i$
- ▶ Process $i = 0$ has a privilege iff $s_0 = s_{N-1}$

Algorithm



- ▶ A privileged process can change its state only via:
 - ▶ For $i > 0$: set $s_i := s_{i-1}$
 - ▶ For $i = 0$: set $s_0 := (s_{N-1} + 1) \bmod K$
- ▶ Note: by this change a process loses its privilege

Recall: Dijkstra's Token Ring - Example

Case 1:

Here $K = 11, N = 5$

Case 2:



Privilege rules (i.e. right to execute critical section):

- ▶ Process $i > 0$ has a privilege iff $s_{i-1} \neq s_i$
- ▶ Process $i = 0$ has a privilege iff $s_0 = s_{N-1}$

Algorithm

- ▶ A privileged process can change its state only via:
 - ▶ For $i > 0$: set $s_i := s_{i-1}$
 - ▶ For $i = 0$: set $s_0 := (s_{N-1} + 1) \bmod K$

Dijkstra's Token Ring: Observations

- ▶ Obs. A: *At least one process has the KS privilege*
 - ▶ If no process $i > 0$ has privilege, then $s_i = s_{i-1}$ for all $i > 0$, so that $s_0 = s_{N-1}$, and thus process 0 has privilege
- ▶ Obs. B: *The number of privileged processes never increases*
 - ▶ The only process that can be privileged is the successor of a privileged process that gives up its privilege

Dijkstra's Token Ring: Observations

- ▶ Obs. A: *At least one process has the KS privilege*
 - ▶ If no process $i > 0$ has privilege, then $s_i = s_{i-1}$ for all $i > 0$, so that $s_0 = s_{N-1}$, and thus process 0 has privilege
- ▶ Obs. B: *The number of privileged processes never increases*
 - ▶ The only process that can be privileged is the successor of a privileged process that gives up its privilege
- ▶ Let the set of legitimate configurations **L** contain the configurations in which exactly one process has the privilege
 - ▶ Configuration here = a union of states of all processes

Lemma: **L** is OK

- ▶ Let the set of legitimate configurations **L** contain the those configurations in which exactly one process has the privilege
 - ▶ Configuration here = a union of states of all processes
- ▶ What is our specification **P** here?
 - ▶ *safety*: only one process is privileged in any configuration (of a configuration sequence)
 - ▶ *liveness*: Any request to become privileged is (eventually) successful
 - ▶ Note: It is essential that P “works” on sequences of config’s!
- ▶ **Lemma**: *Executions that start in a legitimate configuration meet the (problem) specification **P***

Lemma: **L** is OK

- ▶ **Lemma:** *Executions that start in a legitimate configuration meet the (problem) specification **P***
- ▶ **Proof :**
 - ▶ In a configuration of **L**, only the (one) privileged process can take a step
 - ▶ It loses the privilege, but - since there is no configuration without privilege - his successor becomes privileged
 - ▶ Thus, there is always at least one privileged process and each process gets the privilege infinitely often (all N configurations)
 - ▶ Assuming that we take steps at all...

Lemma: Dijkstra's Token Ring Converges

- ▶ **Lemma:** *The algorithm converges to **L** (i.e. each execution contains a configuration of **L**)*
 - ▶ So we always reach to a state where just one process is privileged.
- ▶ **Proof:**
 1. Process 0 performs infinite many steps:
 - A. There can be at most $N(N - 1)/2$ steps without process 0:
 - Consider the function $F(\text{config}) = \sum_{i \in S} (N - i)$
 - where $S = \{i : i > 0 \text{ and } i \text{ is privileged}\}$
 - F 's argument is the current configuration
 - B. F becomes *smaller* with each step of a process $i > 0$, as the privilege goes from i to $i + 1$
 - C. Obviously $F(\text{config}) < N(N - 1)/2$
 - From B. and C. we follow A.

Dijkstra's Token Ring Converges /2

► Proof (2):

2. Process 0 reaches a state w which did not occur in the initial configuration g_0 after at most N of (its own) steps:
 - In g_0 there are at most N different values of states (different numbers among s_0, s_1, \dots, s_{N-1}), so there are still $K - N > 0$ other possible states “outside” g_0
 - Process 0 increases its state modulo K at each of its steps, so it will sometime reach a state w not in g_0 (a bit more thinking required...)
3. When process 0 reaches such a state w for the first time, w does not occur as a state of any other process in the ring:
 - Since other processes only copy states, they can only have states from g_0

Dijkstra's Token Ring Converges /3

► Proof (3):

4. When process 0 gets the privilege the next time, the configuration satisfies $w = s_0 = s_1 = \dots = s_{N-1}$ and is thus in **L**
 - While process 0 has reached the state **w** for the first time, it has lost a privilege and process 1 obtained it
 - Then process 1 lost own privilege by copying value **w** from process 0
 - Then process 2 lost own privilege by copying value **w** from process 1
 - ...
 - Finally, process $N - 1$ got the value **w** from process $N - 2$
 - Now $s_0 = s_{N-1}$ and process 0 is privileged again
 - Note that we have only copied since process 0 has reached the state **w**, so statement #4 is correct!

Dijkstra's Token Ring Converges /4


► Proof (4):

5. This will allow us to achieve a legitimate configuration after at most $N + 1$ steps of process 0
6. Each step of process 0 occurs after at most $N(N - 1)/2$ steps "primitive" steps (statement #1), so the number of steps to stabilization is $O(N^3)$
7. A more detailed analysis can show that this runtime is only $O(N^2)$

Dijkstra's Token Ring is not Uniform

- ▶ In this algorithm, the processes are not the same
- ▶ A stabilizing algorithm is called **uniform** if all processes are the same and have no identity
 - ▶ Uniform solutions are attractive, e.g. you can then equip all with the same "EPROM"
- ▶ Dijkstra showed in 1982 that there is no (uniform) solution to the token ring mutual exclusion problem if the size of the ring is not prime
 - ▶ A uniform solution for ring size = prime p (p is known to all processes) was then found by Burns & Pachl in 1988

Centralized Algorithm for Mutual Exclusion



- Process 1 asks the coordinator for permission to access a shared resource (granted)
- Process 2 then asks for permission to access the same resource; the coordinator does not answer.
- When process 1 releases the resource, it informs the coordinator, who then answers the request of 2

Algorithm of Ricart und Agrawala

- ▶ Ricart and Agrawala have found in 1981 an analogous decentralized algorithm for this problem
- ▶ Idea:
 - ▶ Processes that want to enter the critical section (CS) send a message via multicasting (with a logical time stamp)
 - ▶ They can not enter the CS until all other processes have responded
- ▶ Every process
 - ▶ Has a Lamport clock
 - ▶ Sends messages of the form <Lamport timestamp, ID>
 - ▶ And is in one of the following states:
 - ▶ **RELEASED**: is outside the CS
 - ▶ **WANTED**: process wants to join the CS
 - ▶ **HELD**: is in the CS

Algorithm of Ricart und Agrawala

On initialization

state := RELEASED;

To enter the critical section

state := WANTED;

Multicast request to all processes;

T := request's timestamp;

Wait until (number of replies received = (N - 1));

state := HELD;

On receipt of a request $\langle T_i, p_i \rangle$ at p_j ($i \neq j$)

if (state = HELD or (state = WANTED and $(T, p_j) < (T_i, p_i)$))

then

 queue request from p_i without replying;

else

 reply immediately to p_i ;

end if

i.e. p_j has sent
request before p_i



To exit the critical section

state := RELEASED;

reply to any queued requests;

Algorithmus – Evaluierung

- ▶ N potential points of failure
 - ▶ That is, *all* processes must work intact for the algorithm to work
- ▶ Much communication needed
- ▶ But shows that a fully distributed algorithm is possible for this problem
- ▶ An improved algorithm was developed by Maekawa in 1985
 - ▶ Use the observation that a process does not have to get permission from all other processes, but only from a subset of the processes of size approx. \sqrt{N}

Thank you.

Additional Slides

Arten von fehlertoleranten Algorithmen

Fehlertolerante Algorithmen

- ▶ In VS kann man korrektens Verhalten nicht so leicht garantieren wie in zentralisierten Systemen
 - ▶ Man kann aber die Eigenschaft des **partiellen Ausfalls** (**partial failure**) zum Positiven nutzen
 - ▶ „Make a bug to a feature“
- ▶ Es ist sehr unwahrscheinlich, dass alle Komponenten gleichzeitig ausfallen (Formel?)
 - ▶ Damit kann man den vollständigen Fehlerausfall mit dem Preis der eventuellen Degradation der Systemleistung verhindern
 - ▶ Degradation: geringere Verarbeitungskapazität oder System verhält sich für eine gewisse Zeit inkorrekt
 - ▶ Das ist in zentralisierten Systemen kaum möglich

Robuste vs. Selbststabilisierende Algorithmen

▶ **Robuste Algorithmen**

- ▶ Garantieren korrektes Verhalten der nicht-ausgefallenen Prozesse trotz der Ausfälle / Fehler anderer Prozesse
- ▶ Zu keiner Zeit verhält sich das Gesamtsystem inkorrekt
- ▶ Sie tolerieren **permanente Fehler** einer begrenzten Teilmenge der Komponenten

▶ **Selbststabilisierende Algorithmen**

- ▶ Tolerieren **ein zeitweilig inkorrektes Verhalten** des Systems (**transiente Fehler**)
- ▶ Mehrheit der Komponenten kann sich zeitweise inkorrekt Verhalten
- ▶ Nach einer Weile nimmt das Gesamtsystem korrektes Verhalten an
 - ▶ Wie ein Stehaufmännchen



Robuste Algorithmen

- ▶ Robuste Algorithmen sind „**pessimistisch**“
 - ▶ jede Empfangene Information ist verdächtig
 - ▶ vor jedem Schritt wird eine Reihe von Checks durchgeführt, um die Korrektheit zu gewährleisten
- ▶ Robuste Algorithmen sind **jederzeit korrekt**
 - ▶ zu jeder Zeit der Ausführung wird Korrektheit garantiert
 - ▶ Preis dafür:
 - ▶ die Anzahl der erlaubten fehlerhaften Prozesse ist beschränkt
 - ▶ ggf. die Fehlermodelle schränken mögliche Fehlerarten ein

Entscheidungsprobleme

- ▶ Robusten Algorithmen versuchen meistens, **Entscheidungsprobleme** zu lösen
 - ▶ Jeder Prozess hat ein Input, und alle versuchen daraus zu einer Entscheidung zu kommen
 - ▶ Jeder korrekte Prozesse soll am Ende einen „Entscheidungswert“ in eine spezielle Variable schreiben
 - ▶ Dieser Entscheidungsprozess (z.B. Veto-Wahl) ist trivial ohne Fehler
 - ▶ Fungieren als Bausteine (Primitives) für andere Algorithmen
- ▶ Anforderungen für die Entscheidungen sind meistens:
 - ▶ **Terminierung** (termination)
 - ▶ **Konsistenz** (consistency)
 - ▶ **Nicht-Trivialität** (non-triviality)

Entscheidungsprobleme /2

▶ **Terminierung:**

- ▶ Jeder korrekte Prozess muss schließlich entscheiden

▶ **Konsistenz:**

- ▶ Eine Relation zwischen den Ergebnissen der Prozesse
- ▶ Z.B. **Konsensproblem**: alle Entscheidungen korrekter Prozesse sind gleich
- ▶ Z.B. **Electionproblem**: einer der korrekten Prozesse ist gewählt („leader“), alle anderen haben „verloren“

▶ **Nicht-Trivialität:**

- ▶ Schliesst Lösungen aus, bei denen die Algorithmen gar nicht miteinander kommunizieren bzw. „fest verdrahtet“ sind
- ▶ Z.B. bei Konsensproblem: verboten, dass alle Prozesse sofort „0“ als den Entscheidungswert schreiben

Fehlermodelle

Herbert! Du tot im
Vorzimmer? So kenne ich
Dich aber gar nicht!

▶ Hierarchie von Fehlerarten:

- ▶ **Initial Tote Prozesse** (**initially dead processes**): Prozess führt keinen Schritt des lokalen Algorithmus aus
- In ▶ **Crash** (**crash model**): führt den lokalen Algorithmus für eine Weile korrekt aus, danach stoppt vollständig
- In ▶ **Byzantinisches Verhalten** (**Byzantine behaviour**): Prozess kann sich beliebig Verhalten (und beliebige Nachrichten aussenden)
- ▶ **Merke: jede Art ist ein Spezialfall der nächsten Art**
 - ▶ Damit ist ein byzantinisch-robuster Algorithmus auch Crash-robust etc.
 - ▶ Umgekehrt: *Nichtexistenz* eines Initial-Tot-robusten Algorithmus impliziert die Nichtexistenz der anderen beiden Arten

Beispiel: Konsensproblem

- ▶ Wir möchten einen Algorithmus haben, der einen (binären) Konsens erreicht und dabei bis zu t Crashes toleriert:
 - ▶ **Terminierung**: in jeder *t -crash fairen Ausführung* (d.h. wo mindestens $N-t$ Prozesse unendlich viele Ereignisse ausführen) entscheiden schließlich alle korrekten Prozesse (d.h. setzen unwiderruflich ihren Entscheidungsregister auf einen Wert 0/1)
 - ▶ **Konsistenz**: Keine zwei korrekten Prozesse entscheiden sich für verschiedene Werte
 - ▶ **Nicht-Trivialität**: Es gibt eine Ausführung, bei der die Entscheidung 0 lautet und eine Ausführung, bei der die Entscheidung 1 lautet

Unser Systemmodell

- ▶ Asynchron
 - ▶ Keine obere Schranke auf die Zeit der Ausführung eines Prozessorschritts
 - ▶ Keine obere Schranke auf die Zeit, die eine Nachricht braucht
 - ▶ Damit: keine synchronisierten Uhren möglich
- ▶ Nachrichtenübermittlungsnetzwerk (point-to-point)
- ▶ Keine Übermittlungsfehler
- ▶ Prozesse arbeiten entweder korrekt oder fallen total aus (**crash model**, kein byzantinisches Verhalten)

Bad News: Unmöglichkeit des Konsens

- ▶ Fisher, Lynch, Paterson (1985): Es gibt keinen deterministischen Konsensalgorithmus in diesem Systemmodell, sogar für $t = 1$.
 - ▶ Journal of the ACM, Vol. 32, No. 2, April 1985
- ▶ D.h. sogar unter recht starken Annahmen (keine Übermittlungsfehler, nur ein fehlerhafter Prozessor, keine Byzantinischen Fehler) ist eine „Abstimmung“ unmöglich!
 - ▶ Desto „mehr unmöglich“ für mehr schwerwiegende Fehlerarten
- ▶ Andererseits: kleine Änderungen des Systemmodells lassen Lösungen zu, z.B. Bracha-Toueg-Algorithmus:
 - ▶ Terminierung wird ersetzt durch „Konvergenz“:
 - ▶ Für jede Anfangskonfiguration:

$$\lim_{k \rightarrow \infty} \Pr [\text{ein korrekter Prozess hat nach } k \text{ Schritten noch nicht entschieden}] = 0$$